

VMware Advanced Customer Engagements (ACE) Team

Quick Start Guide for Tanzu Kubernetes Grid Integrated (TKGI) Cluster Backup and Restore

July 2020

Table of Contents

Introduction	3
Recoverability in Kubernetes	3
Velero	4
Velero vSphere plugin	6
Use-case	7
Assumptions	8
Minio	9
Setup Minio.....	9
Minio Cleanup	18
Velero	19
Velero Setup.....	19
Install Velero.....	20
Uninstall Velero	30
Yelb Application	32
Deploy Yelb	32
Unistall Yelb	40
On-demand Backups	41
Cluster Backup Using the Restic Plugin.....	49
Cluster Backup Using the vSphere plugin.....	54
Namespace Backup Using the Restic Plugin	60
Namespace Backup Using the vSphere plugin	62
Scheduled Backups	65
Schedule Backup Using the Restic Plugin.....	65
Schedule Backup Using the vSphere plugin	67
Restore Backups	71
Restore Namespace Backup.....	75
Restore Cluster Backup.....	78
Restore a Point in Time Backup (Scheduled Backup).....	82
Cleanup	86
Conclusion	88

Introduction

This document is a quick start guide for backing up a Tanzu Kubernetes Grid Integrated (TKGI, formerly known as Enterprise PKS) Kubernetes cluster and restoring it. This document will provide details on Valero backup software, installing Velero, backing up an existing cluster, and restoring to the same or another target cluster. The cluster backup will include all Kubernetes (K8) resources as well as persistent volumes.

Recoverability in Kubernetes

Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services at scale and provides high availability that is declarative configuration and automation.

While Kuberentes provides high availability and zero downtime deployments, it does not provide data protection or migration solutions. We need to have a backup and recovery strategy that accounts for:

- Human errors.
- Bugs or vulnerable systems
- Natural disasters
- Legal Requirements for data retention
- Ability to migrate workloads from one cluster to another
 - To deal with changes with providers
 - Cost advantages
 - Compliance with standards.
- Ability to Archive data

We need to have a backup and recovery strategy that will help with the recoverability of Kubernetes. In this guide we will use Velero as the tool but there are several tools available that can be used to accomplish the same objective,

Velero

Running on Kubernetes clusters or on VMs, Velero gives you tools to back up and restore your Kubernetes cluster resources and persistent volumes. You can run Velero in Kubernetes clusters in a public cloud or on-premises. Velero lets you:

- Take backups of your cluster and restore in case of loss.
- Migrate cluster resources to other clusters.
- Replicate your production cluster to development and test clusters.

Velero consists of:

- A server that runs on your cluster
- A command-line client that runs locally

Disaster Recovery

If you periodically back up your cluster's resources, you can return to a previous state in case of some unexpected mishap, such as a service outage.

Cluster Migration

Velero can help you port your resources from one cluster to another, if you point each Velero instance to the same cloud object storage location.

Backup Reference

It is possible to exclude individual items from being backed up, even if they match the resource/namespace/label selectors defined in the backup spec.

Restore Reference

Velero can restore resources into a different namespace than the one they were backed up from.

How it Works

On-demand backups

- Uploads a tar ball of copied Kubernetes objects into cloud object storage.
- Calls the cloud provider API to make disk snapshots of persistent volumes, if specified.

Scheduled backups

- The **schedule** operation allows you to back up your data at recurring intervals

Restores

The **restore** operation allows you to restore all the objects and persistent volumes from a previously created backup. You can also restore only a filtered subset of objects and persistent volumes.

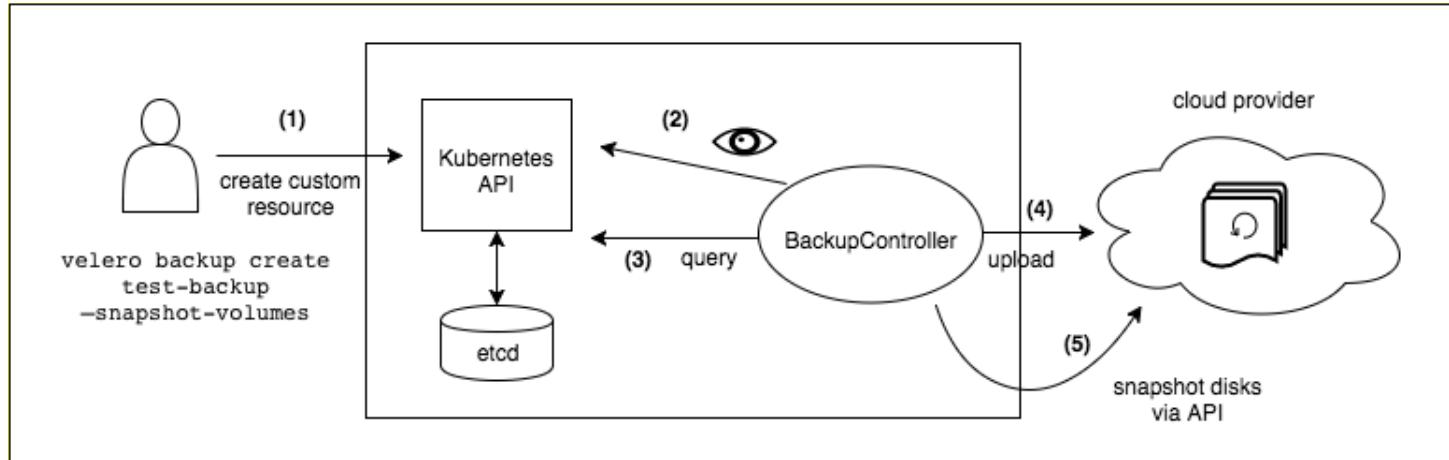
Backup workflow

When you run Velero backup,

- The Velero client makes a call to the Kubernetes API server to create a Backup object.
- The BackupController notices the new Backup object and performs validation.
- The BackupController begins the backup process. It collects the data to back up by querying the API server for resources.

The BackupController makes a call to the object storage service -- for example, AWS (Amazon Web Services) S3 -- to upload the backup.

By default, Velero backup create makes disk snapshots of any persistent volumes. You can adjust the snapshots by specifying additional flags. Run `Velero backup create --help` to see available flags. Snapshots can be disabled with the option `--snapshot-volumes=false`.



You can run Velero in Kubernetes clusters deployed in a public cloud provider or on-premises. For detailed information, see Compatible Storage Providers. Each Velero operation -- on-demand backup, scheduled backup, restore -- is a custom resource, defined with a Kubernetes Custom Resource Definition (CRD) and stored in etcd. Velero also includes controllers that process the custom resources to perform backup, restore, and all related operations. You can back up or restore all objects in your cluster, or you can filter objects by type, namespace, and/or label.

Restic inherently is a file-based backup. Currently, on a vSphere environment Velero uses Restic to backup Kubernetes Persistent Volumes (PV's) by taking the backup of all the files.

For more information go to <https://velero.io/docs/master/restic/>

Velero vSphere plugin

The Velero vSphere plugin enables Velero to take a crash-consistent VMware vSphere Storage APIs snapshot backup of a block Persistent Volume on vSphere storage, and store the backup on S3 compatible storage.

Use-case

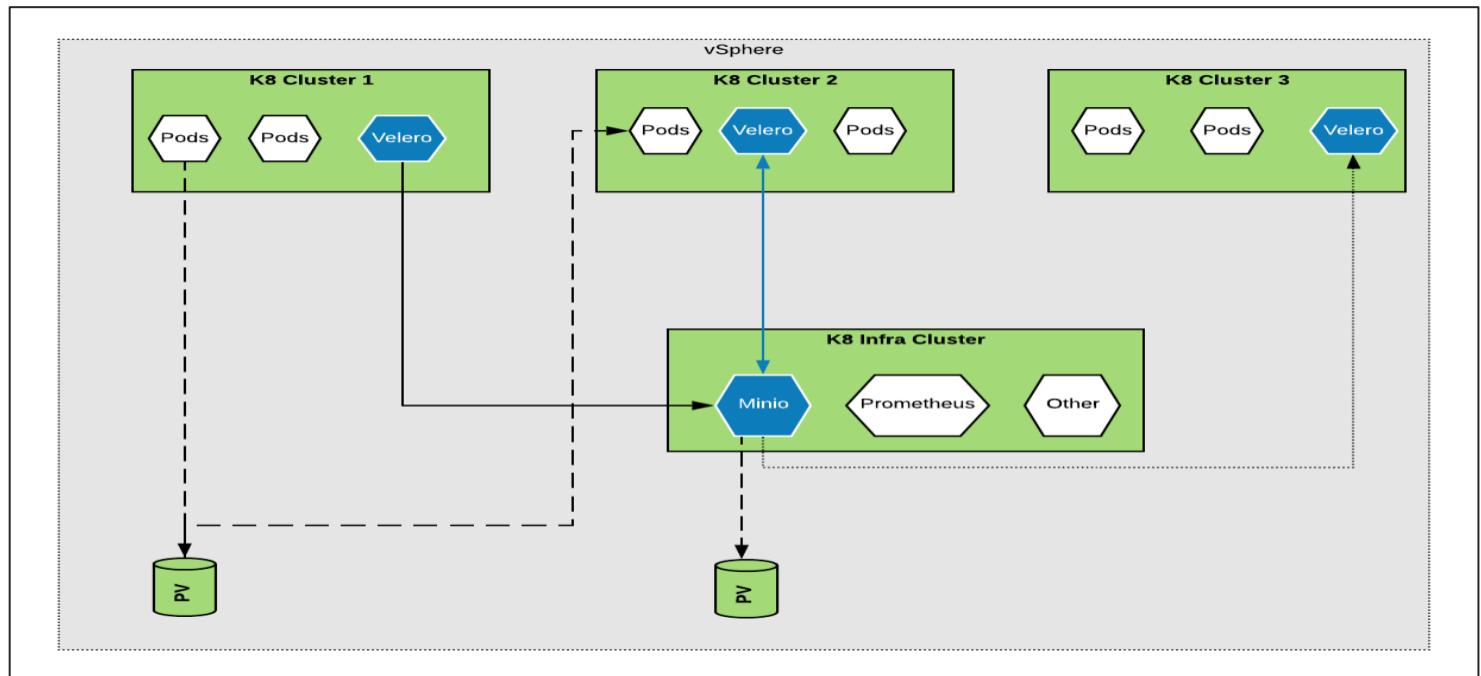
Stateless Applications : these applications do not store data or application state to the cluster or to persistent storage. Instead, data and application state stay with the client, which makes stateless applications more scalable. Those types of applications constitute about 70% of workloads deployed to Kubernetes today, E.g. Front-end applications.

Stateful Applications: these applications, for example databases, save data between sessions and require persistent storage to store the data. The retained data is called the application's state. You can later retrieve the data and use it in the next session. Kubernetes offers persistent volumes as objects capable of retaining their state and data. In the vSphere environment, the persistent volume objects are backed by virtual disks that reside on a datastore.

This section defines the common use cases where Velero would be applicable.

Cluster Migration: Migrate both stateful and stateless applications from one cluster to another

Disaster Recovery: Restore applications (both stateful and stateless) to a cluster from a backup in time (scheduled backup)



Assumptions

The following assumptions are made in the guide:

- TKGI is deployed
- The infrastructure team has setup 3 K8s clusters, the source cluster (where Velero backup is made from) , the target (where the Velero backup is restored to) cluster and an infra cluster where all infrastructure applications like Minio and Prometheus will be running .

NOTE: This is not a fixed rule, Minio can run on any cluster, including the source and the target cluster, it could also be run on a standalone vm. For more information on Minio visit <https://docs.min.io/>

- The Minio backup endpoint is accessible from both the source and target clusters.
- A Linux/ubuntu machine is provisioned to install and use various software components
- The provisioned Linux/ubuntu machine meets the following
 - Can access all the 3 K8s clusters defined above
 - Has the appropriate kubectl cli installed
 - Has the appropriate pks cli installed
 - Has the latest version of *Helm* installed?
- In this document, we will be using ci-cluster as our source cluster and my-cluster as our target cluster.

Minio

Setup Minio

We will be setting up an open source version of Minio in the infrastructure cluster. We will be using the VMware Bitnami official opensource Minio Helm chart for K8s. The steps below describe how to setup Minio in a K8 cluster using the Bitnami distribution.

Step 1: ssh to the provisioned ubuntu vm (clivm)

Step 2: Get kube config for the infra cluster

```
pks login -a <pks api> -u <user> -p <password> -k  
pks get-credentials <cluster>
```

Alternatively

```
pks get-kubeconfig <cluster> -a <pks api> -u <user> -p <password> -k
```

E.g.

```
pks login -a pks.corp.local -u riaz -p VMware1! -k
```

```
pks get-credentials infra-cluster
```

Or

```
pks get-kubeconfig infra-cluster -a pks.corp.local -u riaz -p VMware1! -k
```

Step 3: Create a namespace to which Minio can be deployed

```
kubectl create ns minio
```

Step 4: Add Bitnami Helm repository

```
helm repo add bitnami https://charts.bitnami.com/bitnami
```

```
helm repo update
```

Step 5: Minio requires a backing store / K8s persistent volume. Create a storage-class on the infra cluster with the following storage class definition. Copy the contents of the file below to a file storage-class.yaml and create the storage class.

```
---
```

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: minio-disk
provisioner: kubernetes.io/vsphere-volume
parameters:
  diskformat: thin
```

NOTE: If setting up the storage class using the CSI driver , follow the steps provided in the VMware Tanzu documentation to set up the CSI driver on the cluster you before creating the storage class.
<https://docs.pivotal.io/pks/1-7/vsphere-cns.html>

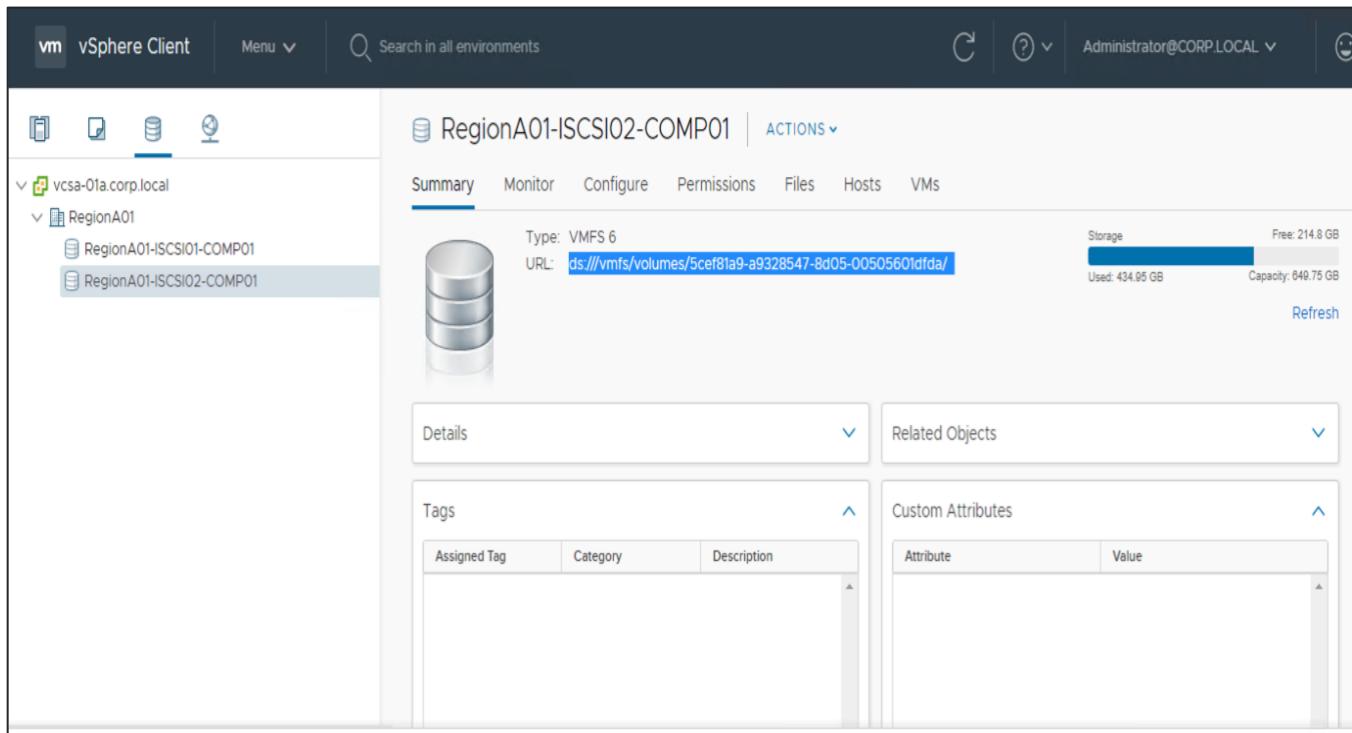
Storage class definition when using a CSI driver

```
---
```

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: minio-disk
annotations:
  storageclass.kubernetes.io/is-default-class: "true"
provisioner: csi.vsphere.vmware.com
parameters:
  datastoreurl: "ds:///vmfs/volumes/5cef81a9-a9328547-8d05-00505601dfda/"
```

TKG Backup and Restore Clusters

Datastore url can be obtained from vCenter



The screenshot shows the vSphere Client interface. On the left, the navigation tree shows a folder named 'RegionA01' containing two items: 'RegionA01-ISCSI01-COMP01' and 'RegionA01-ISCSI02-COMP01'. The 'RegionA01-ISCSI02-COMP01' item is selected. The main pane displays the 'Summary' tab for this object. The 'Type' is listed as 'VMFS 6'. The 'URL' field contains the value 'ds://vmfs/volumes/5cef81a9-a9328547-8d05-00505601dtda/'. Below this, there are sections for 'Details' (which is expanded) and 'Related Objects'. The 'Tags' section is empty. The 'Custom Attributes' section is also empty. The top right corner shows the user 'Administrator@CORP.LOCAL'.

```
kubectl apply -f storage-class.yaml
```

Step 6: Deploy the Bitnami Minio release. This will create the necessary resources to run Minio within the minio namespace

```
helm install minio-release -n minio --set accessKey.password=minio --set  
secretKey.password=minio123 --set persistence.storageClass=minio-disk --set persistence.size=128Gi  
bitnami/minio
```

NOTE: When sizing the storage for Minio, make sure there is about 3X times the storage the reason for which is minio stages the backup before commit. E.g. if a backup takes about 10Gi we would require about 30Gi .,. The extra storage is for the restore , velero uploads restore data to minio which also requires space.

Step 7: Check for all pods, deployments and services and make sure everything is created and the pods are running as expected. Also check if the PVC is created and bound

```
kubectl get all -n minio  
kubectl get pvc -n minio  
kubectl get deployment -n minio
```

```
ubuntu@cli-vm:~/velero$ kubectl get all -n minio  
NAME                                     READY   STATUS    RESTARTS   AGE  
pod/minio-release-d8b746dd8-lbsrw        1/1     Running   0          5m26s  
  
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE  
service/minio-release   ClusterIP  10.100.200.107  <none>        9000/TCP   5m27s  
  
NAME           READY   UP-TO-DATE   AVAILABLE   AGE  
deployment.apps/minio-release   1/1     1           1          5m27s  
  
NAME           DESIRED  CURRENT  READY   AGE  
replicaset.apps/minio-release-d8b746dd8  1        1        1       5m27s
```

```
ubuntu@cli-vm:~/velero$ kubectl get pvc -n minio  
NAME      STATUS  VOLUME                                     CAPACITY  ACCESS MODES  STORAGECLASS  AGE  
minio-release  Bound  pvc-66b6da36-e06e-4b0b-96c5-67efc3f2a1f8  8Gi      RWO          minio-disk  5m59s
```

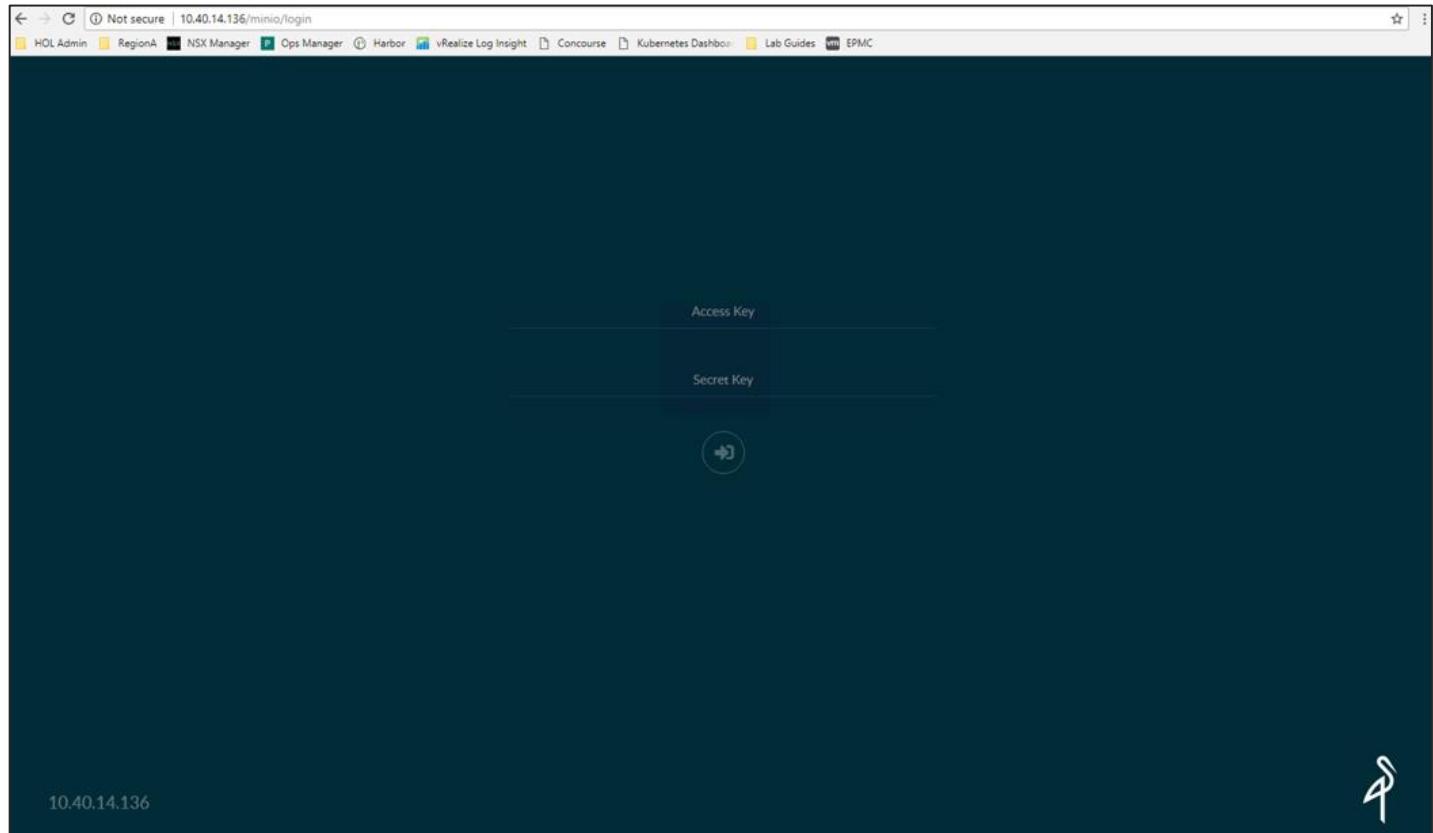
```
ubuntu@cli-vm:~/velero$ kubectl get deployment -n minio  
NAME           READY   UP-TO-DATE   AVAILABLE   AGE  
minio-release  1/1     1           1          9m3s
```

Step 8: Expose the deployment as a Load Balancer. This will create a lb within NSX-T as well as an ingress.

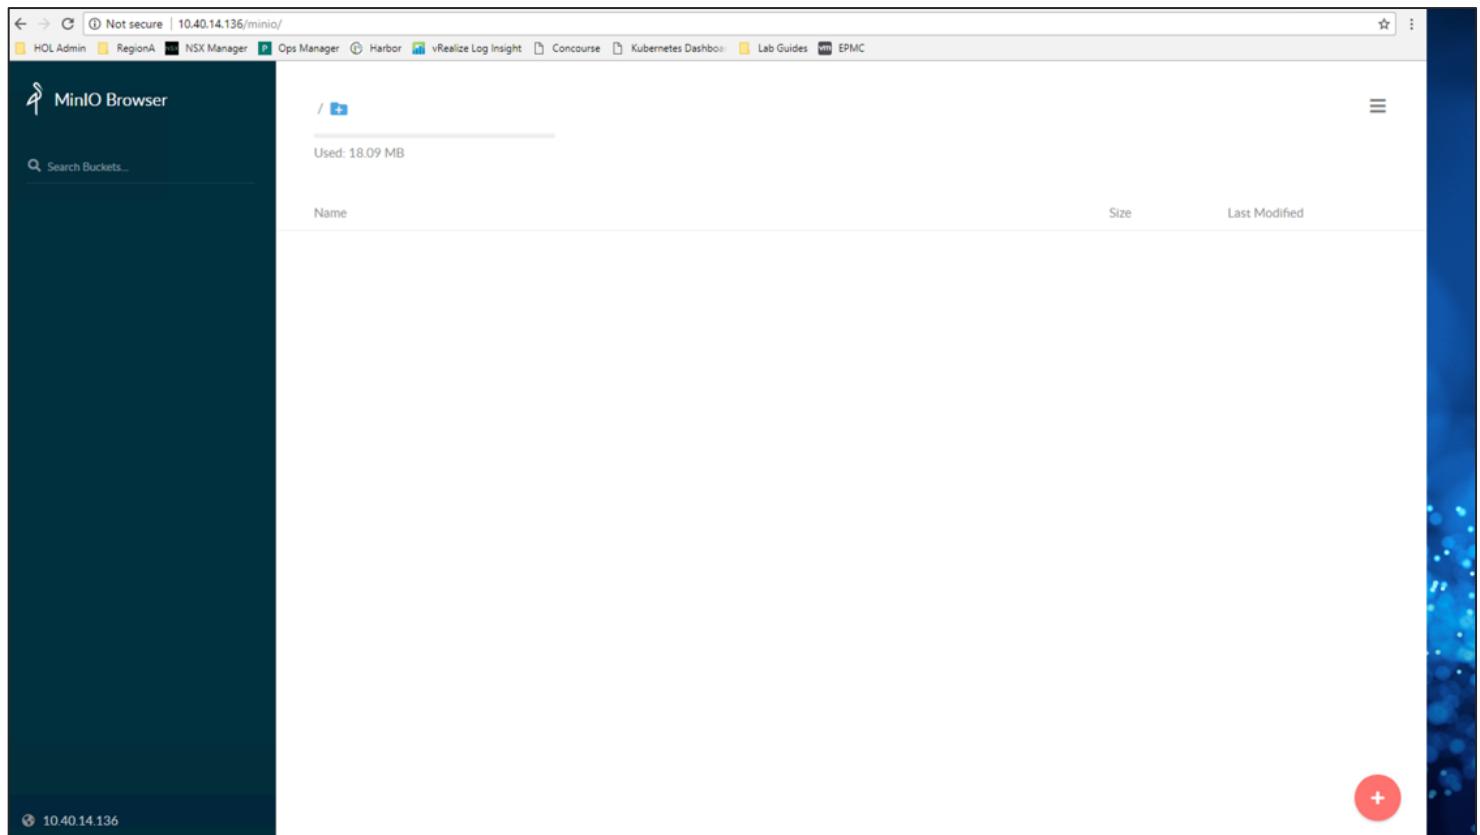
```
kubectl expose deployment minio-release --name=minio-frontend-lb --port=80 --target-port=9000 --  
type=LoadBalancer --namespace=minio
```

Step 9: Check the IP under the “External-IP” section, point your browser to that IP address:port . The Minio application should be accessible

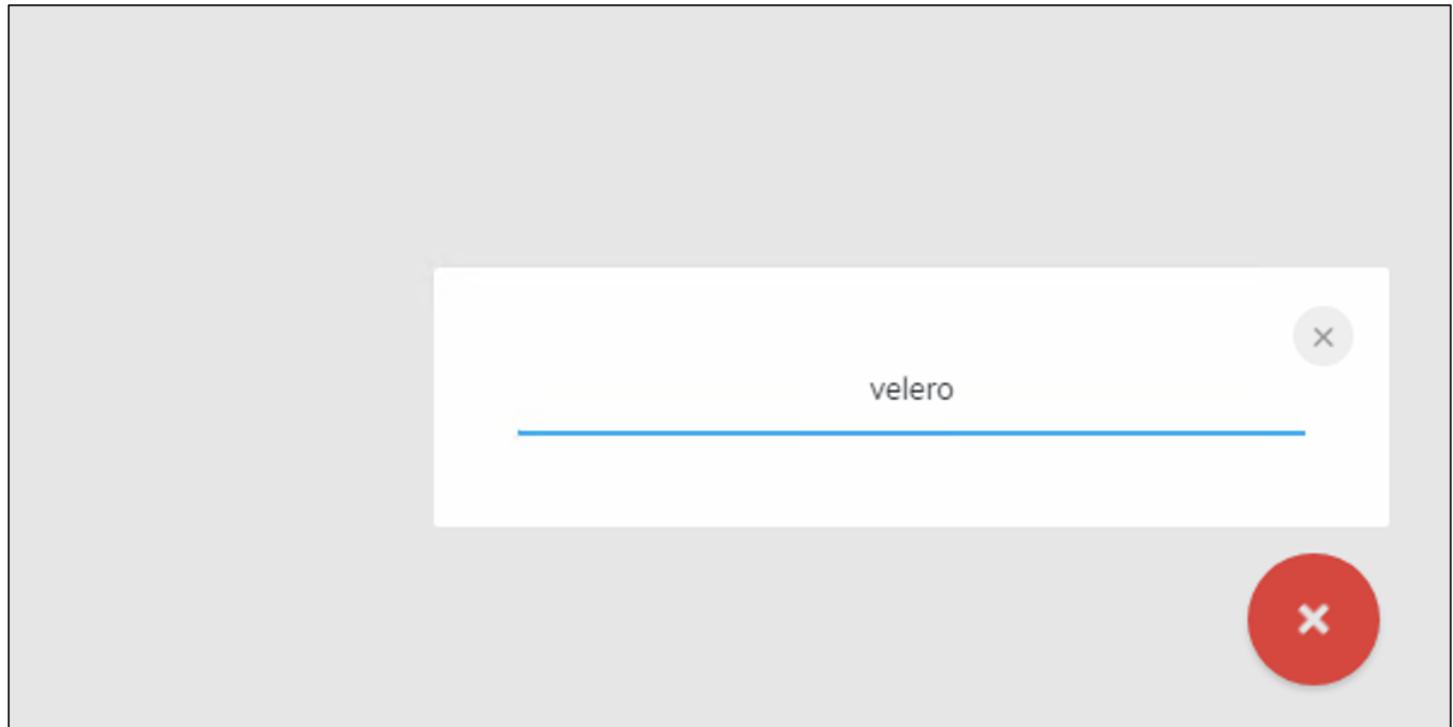
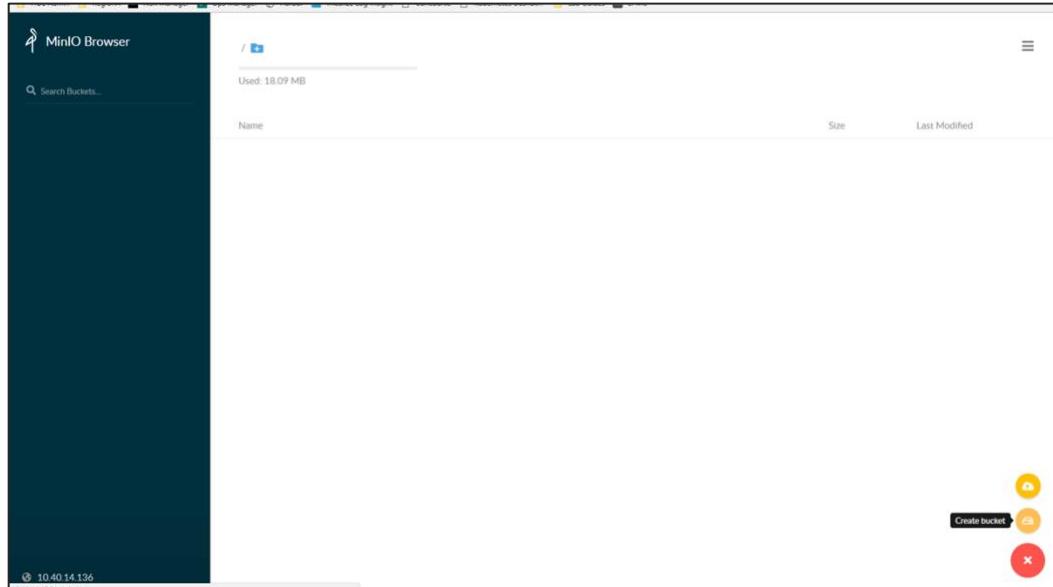
```
ubuntu@cli-vm:~/velero$ kubectl get svc -n minio
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP     PORT(S)        AGE
minio-frontend-lb   LoadBalancer   10.100.200.3  10.40.14.136  80:30568/TCP  18s
minio-release     ClusterIP    10.100.200.107 <none>        9000/TCP     13m
```

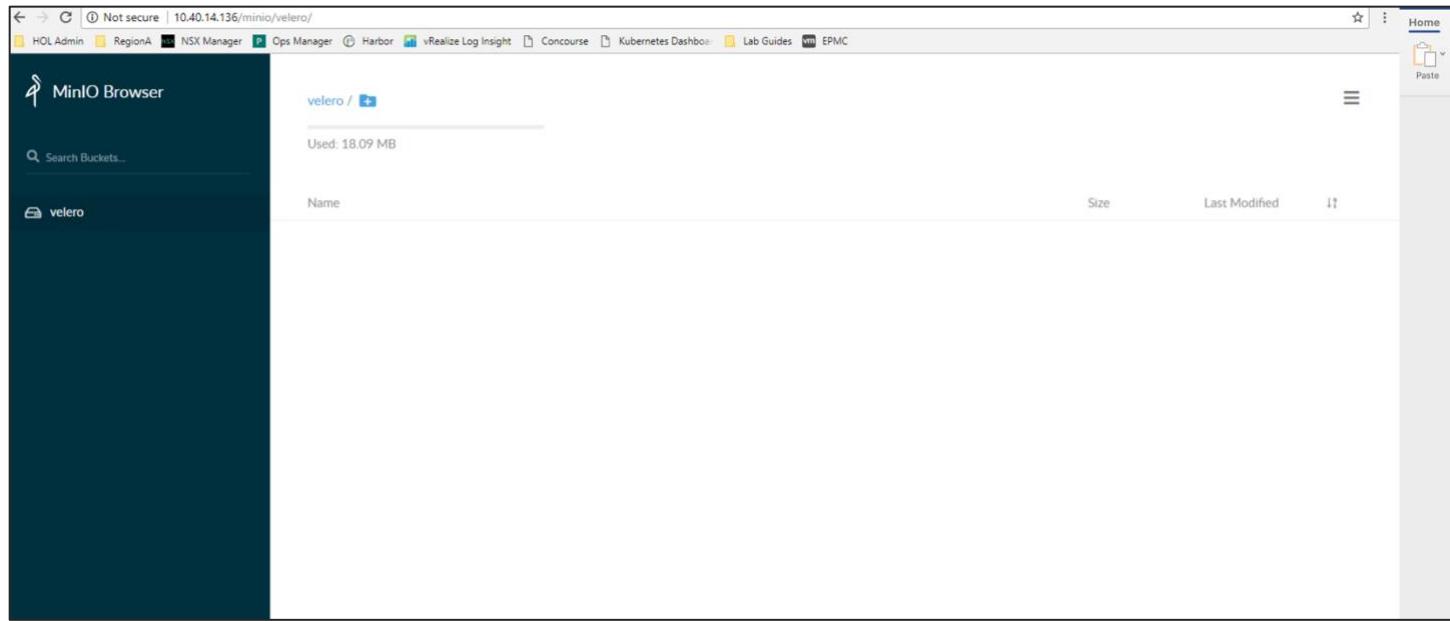


Step 10: Login with the credentials used in step 6. E.g. minio/minio123



Step 11: Create a bucket called Velero. We will be using this bucket when we install Velero to the clusters in the following steps:





Minio Cleanup

To clean up Minio, uninstall the deployed helm release

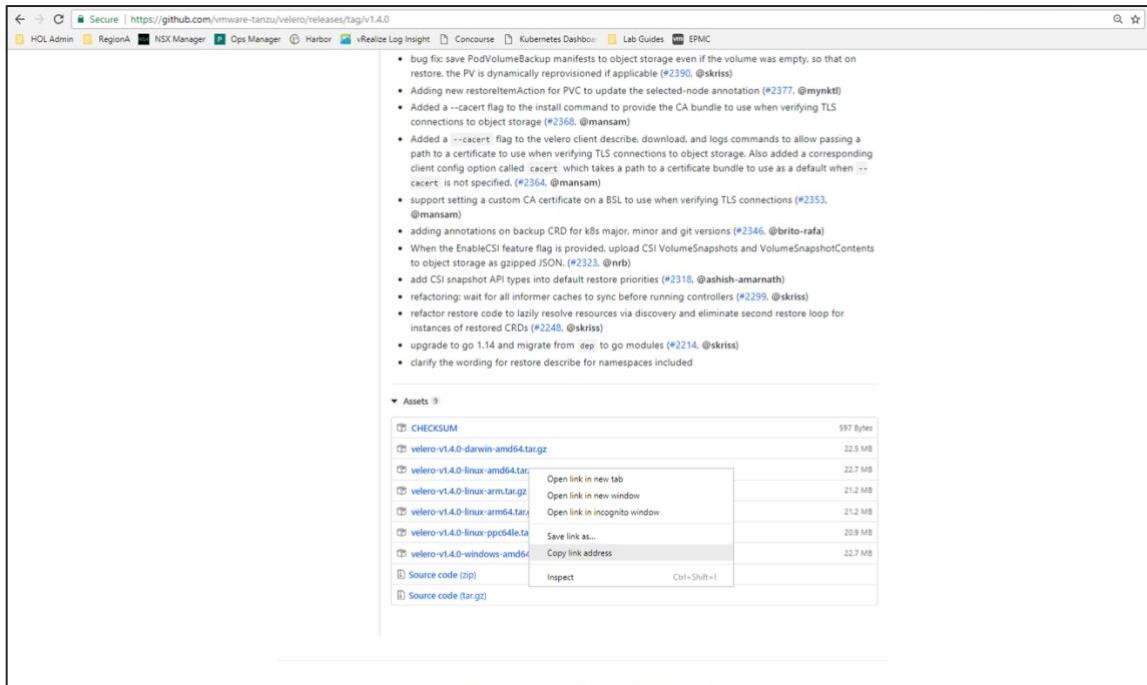
```
helm ininstall minio-release -n minio
```

Velero

This section goes through the steps to download Velero to the provisioned Ubuntu vm (cli vm) and install it on both the source and target clusters.

Velero Setup

Step 1: Navigate to the official page of Velero (<https://github.com/vmware-tanzu/velero/releases>) and copy the link for the target VM (Virtual Machine) OS (operating systems). (Eg. <https://github.com/vmware-tanzu/velero/releases/tag/v1.4.0>). At the bottom of the page the official releases are listed, Right click on the release link 'Copy Link address'



Step 2: ssh to the provisioned ubuntu vm. (clivm)

Step 3: Download and uncompress the Velero distribution

```
mkdir velero  
cd ~/velero  
wget https://github.com/vmware-tanzu/velero/releases/download/v1.4.0/velero-v1.4.0-linux-amd64.tar.gz  
tar xvf velero-v1.4.0-linux-amd64.tar.gz
```

Install Velero

This section describes the steps required to install Velero to both the source and target clusters. Any cluster from which a backup is taken or to which a backup is restored requires to have Velero deployed to it.

Source Cluster

Source cluster is the cluster from which a Velero backup will be taken from. As mentioned in the assumptions section we will be using the ci-cluster as our source cluster.

Step 1: ssh into the provisioned linux/ubuntu vm (clivm)

Step 2: Get kube config for the source cluster:

```
pks login -a <pks api> -u <user> -p <password> -k  
pks get-credentials <cluster>
```

Alternatively

```
pks get-kubeconfig <cluster> -a <pks api> -u <user> -p <password> -k
```

E.g.



```
pks login -a pks.corp.local -u riaz -p VMware1! -k  
pks get-credentials ci-cluster  
pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k
```

```
ubuntu@cli-vm:~/velero$ pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k  
Fetching kubeconfig for cluster ci-cluster and user riaz.  
You can now use the kubeconfig for user riaz:  
$kubectl config use-context ci-cluster
```

Step 3: Create a namespace for Velero, called velero:

```
kubectl create ns velero
```

Step 4: Change directory to the velero directory:

```
cd ~/velero/velero-v1.4.0-linux-amd64
```

Step 5: Create a credentials file, and name it credentials. This will contain the username and password used for Minio. The values would be the same as what was provided during the Minio setup.

```
[default]  
aws_access_key_id = minio  
aws_secret_access_key = minio123
```

NOTE: This file can be deleted once the Velero is installed to the cluster.

Step 6: Set kubectl context to the source cluster

```
kubectl config use-context <source-cluster>
```

E.g.

```
kubectl config use-context ci-cluster
```

Step 7: Install Velero to the source cluster.

```
./velero install \
--provider aws \
--bucket velero \
--secret-file ./credentials \
--use-restic \
--backup-location-config region=minio,s3ForcePathStyle="true",s3Url=http://<external-ip of
minio>:<port> \
--snapshot-location-config region=minio \
--plugins velero/velero-plugin-for-aws:v1.1.0
```

Note: The secrets file points to location of the file credentials file we created above
use Restic to backup pv's the s3Url points to the Minio that was setup earlier.

E.g.

```
./velero install \
--provider aws \
--bucket velero \
--secret-file ./credentials \
--use-restic \
--backup-location-config region=minio,s3ForcePathStyle="true",s3Url=http://10.40.14.66 \
--snapshot-location-config region=minio \
--plugins velero/velero-plugin-for-aws:v1.1.0
```

TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero$ ./velero install \
>   --provider aws \
>   --bucket velero \
>   --secret-file ./credentials \
>   --backup-location-config region=minio,s3ForcePathStyle="true",s3Url=http://10.40.14.66 \
>   --snapshot-location-config region=minio \
>   --plugins velero/velero-plugin-for-aws:v1.1.0
CustomResourceDefinition/backups.velero.io: attempting to create resource
CustomResourceDefinition/backups.velero.io: created
CustomResourceDefinition/backupstoragelocations.velero.io: attempting to create resource
CustomResourceDefinition/backupstoragelocations.velero.io: created
CustomResourceDefinition/deletebackuprequests.velero.io: attempting to create resource
CustomResourceDefinition/deletebackuprequests.velero.io: created
CustomResourceDefinition/downloadrequests.velero.io: attempting to create resource
CustomResourceDefinition/downloadrequests.velero.io: created
CustomResourceDefinition/podvolumebackups.velero.io: attempting to create resource
CustomResourceDefinition/podvolumebackups.velero.io: created
CustomResourceDefinition/podvolumerestores.velero.io: attempting to create resource
CustomResourceDefinition/podvolumerestores.velero.io: created
CustomResourceDefinition/restrictrepositories.velero.io: attempting to create resource
CustomResourceDefinition/restrictrepositories.velero.io: created
CustomResourceDefinition/restores.velero.io: attempting to create resource
CustomResourceDefinition/restores.velero.io: created
CustomResourceDefinition/schedules.velero.io: attempting to create resource
CustomResourceDefinition/schedules.velero.io: created
CustomResourceDefinition/serverstatusrequests.velero.io: attempting to create resource
CustomResourceDefinition/serverstatusrequests.velero.io: created
CustomResourceDefinition/volumesnapshotlocations.velero.io: attempting to create resource
CustomResourceDefinition/volumesnapshotlocations.velero.io: created
Waiting for resources to be ready in cluster...
Namespace/velero: attempting to create resource
Namespace/velero: created
ClusterRoleBinding/velero: attempting to create resource
ClusterRoleBinding/velero: created
ServiceAccount/velero: attempting to create resource
ServiceAccount/velero: created
Secret/cloud-credentials: attempting to create resource
Secret/cloud-credentials: created
BackupStorageLocation/default: attempting to create resource
BackupStorageLocation/default: created
VolumeSnapshotLocation/default: attempting to create resource
VolumeSnapshotLocation/default: created
Deployment/velero: attempting to create resource
Deployment/velero: created
Velero is installed! Use 'kubectl logs deployment/velero' to view the status.
```

Step 8: Get status of pods in the velero namespace

```
kubectl get po -n velero
```

Step 9: If the Restic pods fails to startup, we will need to edit the hostPath for the Restic pods.

```
kubectl edit daemonset restic -n velero
```

change hostPath from /var/lib/kubelet/pods to /var/vcap/data/kubelet/pods:

Which will look like below

```
-hostPath:
  path: /var/vcap/data/kubelet/pods
```

VMware®

TKG Backup and Restore Clusters

```
- name: VELERO_NAMESPACE
  valueFrom:
    fieldRef:
      apiVersion: v1
      fieldPath: metadata.namespace
- name: VELERO_SCRATCH_DIR
  value: /scratch
- name: GOOGLE_APPLICATION_CREDENTIALS
  value: /credentials/cloud
- name: AWS_SHARED_CREDENTIALS_FILE
  value: /credentials/cloud
- name: AZURE_CREDENTIALS_FILE
  value: /credentials/cloud
- name: ALIBABA_CLOUD_CREDENTIALS_FILE
  value: /credentials/cloud
image: velero/velero:v1.4.0
imagePullPolicy: IfNotPresent
name: restic
resources: {}
terminationMessagePath: /dev/termination-log
terminationMessagePolicy: File
volumeMounts:
- mountPath: /host_pods
  mountPropagation: HostToContainer
  name: host-pods
- mountPath: /scratch
  name: scratch
- mountPath: /credentials
  name: cloud-credentials
dnsPolicy: ClusterFirst
restartPolicy: Always
schedulerName: default-scheduler
securityContext:
  runAsUser: 0
serviceAccount: velero
serviceAccountName: velero
terminationGracePeriodSeconds: 30
volumes:
- hostPath:
    path: /var/vcap/data/kubelet/pods
    type: ""
  name: host-pods
- emptyDir: {}
  name: scratch
- name: cloud-credentials
  secret:
    defaultMode: 420
    secretName: cloud-credentials
templateGeneration: 1
updateStrategy:
  rollingUpdate:
    maxUnavailable: 1
  type: RollingUpdate
status:
  currentNumberScheduled: 3
  desiredNumberScheduled: 3
  numberMisscheduled: 0
  numberReady: 0
  numberUnavailable: 3
  observedGeneration: 1
  updatedNumberScheduled: 3
- INSERT --
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get po -n velero
NAME          READY   STATUS    RESTARTS   AGE
restic-hmzfl  1/1     Running   0          43m
restic-jff8c  1/1     Running   0          43m
restic-s9flx  1/1     Running   0          43m
velero-84d944c59-2c9rv 1/1     Running   0          94s
```

Step 10: Get all the plugins in velero

```
./velero plugin get
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero plugin get
NAME                      KIND
velero.io/crd-remap-version BackupItemAction
velero.io/pod                BackupItemAction
velero.io/pv                 BackupItemAction
velero.io/service-account    BackupItemAction
velero.io/aws                ObjectStore
velero.io/add-pv-from-pvc   RestoreItemAction
velero.io/add-pvc-from-pod  RestoreItemAction
velero.io/change-pvc-node-selector RestoreItemAction
velero.io/change-storage-class RestoreItemAction
velero.io/cluster-role-bindings RestoreItemAction
velero.io/crd-preserve-fields RestoreItemAction
velero.io/job                RestoreItemAction
velero.io/pod                RestoreItemAction
velero.io/restic             RestoreItemAction
velero.io/role-bindings     RestoreItemAction
velero.io/service            RestoreItemAction
velero.io/service-account   RestoreItemAction
velero.io/aws                VolumeSnapshotter
```

Note: The AWS and restic plugins installed in the previous step is listed.

Step 11: Install the Velero Plugin for vSphere. This plugin is an alternate to Restic and enables Velero to take crash-consistent snapshot backup of a block Persistent Volume on vSphere storage and backup of volume data into S3 compatible storage. The Velero vSphere plugin has a few prerequisites:

- Velero - Version 1.3.2 or above
- vSphere - Version 6.7U3 or above
- vSphere CSI/CNS driver 1.0.2 or above
- Kubernetes 1.14 or above (note: the Velero Plug-in for vSphere does not support Guest or Supervisor clusters on vSphere yet)

Make sure these are met before you proceed to the next step

Step 12: Install the vSphere CSI driver . Follow the steps provided in the VMware Tanzu documentation to set up the CSI driver on the cluster you before creating the storage class. <https://docs.pivotal.io/pks/1-7/vsphere-cns.html>

NOTE: Make sure that the csi-vsphere.conf file has the correct values

[Global]

cluster-id = <cluster-name>

[VirtualCenter "<vcenter-ip>"]

insecure-flag = "true"

user = "<vcenter-username>"

password = "vcenter-password"

port = "443"

datacenters = "<vcenter-datacenter>"

Eg.

[Global]

cluster-id = ci-cluster

[VirtualCenter "192.168.110.22"]

insecure-flag = "true"

user = "administrator@corp.local"

password = "VMware1!"

```
port = "443"  
datacenters = "RegionA01"
```

Step 13: Install the velero vsphere plugin

```
./velero plugin add vsphereveleroplugin/velero-plugin-for-vsphere:1.0.1
```

Step 14: Verify if the velero vsphere plugin has been installed

```
./velero plugin get
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero plugin get
NAME          KIND
velero.io/crd-remap-version   BackupItemAction
velero.io/pod                 BackupItemAction
velero.io/pv                  BackupItemAction
velero.io/service-account     BackupItemAction
velero.io/aws                 ObjectStore
velero.io/add-pv-from-pvc    RestoreItemAction
velero.io/add-pvc-from-pod   RestoreItemAction
velero.io/change-pvc-node-selector RestoreItemAction
velero.io/change-storage-class RestoreItemAction
velero.io/cluster-role-bindings RestoreItemAction
velero.io/crd-preserve-fields RestoreItemAction
velero.io/job                 RestoreItemAction
velero.io/pod                 RestoreItemAction
velero.io/restic              RestoreItemAction
velero.io/role-bindings       RestoreItemAction
velero.io/service             RestoreItemAction
velero.io/service-account     RestoreItemAction
velero.io/aws                 VolumeSnapshotter
velero.io/vsphere              VolumeSnapshotter
```

Note : The vsphere plugin is listed along with AWS and restic.

Target Cluster

A target cluster is the cluster to which a Velero backup is to be restored. As mentioned in the assumptions section we will be using ‘my-cluster’ as our source cluster

Step 1: ssh into the provisioned linux/ubuntu vm (clivm)

Step 2: Get kube config for the target cluster

```
pks login -a <pks api> -u <user> -p <password> -k  
pks get-credentials <cluster>
```

Alternatively

```
pks get-kubeconfig <cluster> -a <pks api> -u <user> -p <password> -k
```

E.g.

```
pks login -a pks.corp.local -u riaz -p VMware1! -k  
pks get-credentials my-cluster  
pks get-kubeconfig my-cluster -a pks.corp.local -u riaz -p VMware1! -k
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ pks get-kubeconfig my-cluster -a pks.corp.local -u riaz -p VMware1! -k  
Fetching kubeconfig for cluster my-cluster and user riaz.  
You can now use the kubeconfig for user riaz:  
$kubectl config use-context my-cluster
```

Step 3: Follow steps 3 to 14 in the previous section.

Uninstall Velero

To uninstall Velero we will need to delete all the resources associated with its install

```
kubectl delete namespace/velero clusterrolebinding/velero
```

```
kubectl delete crds -l component=velero
```

Yelb Application

In this section we will be deploying the Yelb (<https://github.com/mreferre/yelb>) application to the source cluster. Yelb allows users to vote on a set of alternatives (restaurants) and dynamically updates pie charts based on number of votes received.

The Yelb application will help us validate stateful pods and persistent data is being backed up by Velero and the Velero restore operation keeps the state of the data intact.

Deploy Yelb

Step 1: Get kube config for the source cluster

```
pks login -a <pks api> -u <user> -p <password> -k  
pks get-credentials <cluster>
```

Alternatively

```
pks get-kubeconfig <cluster> -a <pks api> -u <user> -p <password> -k
```

E.g.

```
pks login -a pks.corp.local -u riaz -p VMware1! -k  
pks get-credentials ci-cluster  
pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k
```

```
ubuntu@cli-vm:~/velero$ pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k  
Fetching kubeconfig for cluster ci-cluster and user riaz.  
You can now use the kubeconfig for user riaz:  
$kubectl config use-context ci-cluster
```

Step 2: Set kubectl context to the source cluster

```
kubectl config use-context <source-cluster>
```

E.g.

```
kubectl config use-context ci-cluster
```

Step 3: Yelb requires a backing store / K8s persistent volume. Create a storage-class on the infra cluster with the following storage class definition. Copy the contents of the file below to a file storage-class.yaml and create the storage class.

```
---
```

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: thin-disk
annotations:
  storageclass.kubernetes.io/is-default-class: "true"
provisioner: kubernetes.io/vsphere-volume
parameters:
  diskformat: thin
```

NOTE: If setting up the storage class using the csi driver , follow the steps provided in the VMware Tanzu documentation to set up the CSI driver on the cluster you before creating the storage class.

<https://docs.pivotal.io/pks/1-7/vsphere-cns.html>

Storage class definition when using a CSI driver

```
---
```

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: csi-sc
annotations:
  storageclass.kubernetes.io/is-default-class: "true"
provisioner: csi.vsphere.vmware.com
parameters:
```

datastoreurl: "ds:///vmfs/volumes/5cef81a9-a9328547-8d05-00505601dfda/"

TKG Backup and Restore Clusters

Datastore url can be obtained from vcenter

The screenshot shows the vSphere Client interface. On the left, the navigation tree shows a folder named 'RegionA01' containing three items: 'RegionA01-ISCSI01-COMP01', 'RegionA01-ISCSI02-COMP01', and 'RegionA01-ISCSI03-COMP01'. The 'RegionA01-ISCSI02-COMP01' item is selected and highlighted in blue. The main pane displays the 'RegionA01-ISCSI02-COMP01' details. The 'Summary' tab is selected. The 'Type' is listed as 'VMFS 6'. The 'URL' is shown as a blue hyperlink: <ds:///vmfs/volumes/5cef81a9-a9328547-8d05-00505601dfa/>. In the 'Storage' section, it shows 'Free: 214.8 GB', 'Used: 434.85 GB', and 'Capacity: 640.75 GB'. A 'Refresh' button is located at the bottom right of the storage section.

kubectl apply -f storage-class.yaml

Step 3: Create a yelb namespace

kubectl create ns yelb

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl create ns yelb
namespace/yelb created
```


Step 4: Create the persistent volume claims required for this application. This will create a PVC for the redis-server and the yelb-db database.

Copy the contents of the file

<https://github.com/riazvm/velerobackupandrestore/blob/master/application/yelb/yelb-pvc.yaml>
to a local file for e.g. yelb-pvc.yaml

```
kubectl apply -f yelb-pvc.yaml
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl apply -f yelb-pvc.yaml
persistentvolumeclaim/redis-pv-claim created
persistentvolumeclaim/db-pv-claim created
```

Step 5: Deploy the Yelb application, this will create the necessary deployments, pods and services and expose the yelb-ui deployment as an ingress of type loadbalancer

Copy the contents of the file

<https://github.com/riazvm/velerobackupandrestore/blob/master/application/yelb/yelb.yaml>
file to a local file for e.g. yelb.yaml

```
kubectl apply -f yelb.yaml
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl apply -f yelb.yaml
service/redis-server created
service/yelb-db created
service/yelb-appserver created
service/yelb-ui created
deployment.extensions/yelb-ui created
statefulset.apps/redis-server created
statefulset.apps/yelb-db created
deployment.extensions/yelb-appserver created
```

Step 6: Verify if all the pods are running on the yelb name namespace

```
kubectl get po -n yelb --watch
```

TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get po -n yelb --watch
NAME                    READY   STATUS    RESTARTS   AGE
redis-server-0          1/1     Running   0          69s
yelb-appserver-696b9668c4-7mrhs 1/1     Running   0          69s
yelb-db-0               1/1     Running   0          69s
yelb-ui-6665575695-2lqzw   1/1     Running   0          68s
```

Step 7: Get the loadbalancer ip for the ingress to access the Yelb application. The EXTERNAL-IP is the loadbalancer's ip.

```
kubectl get svc -n yelb
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get svc -n yelb
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
redis-server  ClusterIP  10.100.200.113 <none>        6379/TCP      2m12s
yelb-appserver  ClusterIP  10.100.200.155 <none>        4567/TCP      2m11s
yelb-db      ClusterIP  10.100.200.158 <none>        5432/TCP      2m12s
yelb-ui      LoadBalancer  10.100.200.3  10.40.14.46  80:30561/TCP  2m11s
```

Step 8: Access the application by pointing the browser to the ip from the previous step

The screenshot shows a web browser window with the URL 10.40.14.46. The page title is "Yelb". The content displays a welcome message: "Welcome to Yelb, the only hub for healthy food recommendations!". Below this, there are four cards representing different restaurants:

- IHOP**: Shows a logo of two eggs, the text "Pancakes, for a powerful start!", and a blue "VOTE" button.
- Chipotle**: Shows a logo of a burrito, the text "Burritos, for a mid-day break!", and a blue "VOTE" button.
- Outback**: Shows a logo of a blooming onion, the text "Blooming onion, what else?", and a blue "VOTE" button.
- Buca di Beppo**: Shows a logo of a lasagna, the text "Lasagne, a great taste directly from Italy", and a blue "VOTE" button.

Below the cards, there is a chart titled "Total" showing the distribution of votes across the four categories. The chart has four bars: Outback (green), Buca di Beppo (red), IHOP (yellow), and Chipotle (blue). The values are all NaN%.

At the bottom left, there is a box stating "Number of page views so far: 1". At the bottom right, there is a box stating "System serving the request: yelb-appserver-696b9668c4-nbjp4".

Unistall Yelb

To delete the Yelb application and all its associated resources, delete the yelb namespace

```
kubectl delete ns yelb
```

On-demand Backups

This section describes steps to backup resources deployed to a source cluster. Before backing up an application , namespace, or a cluster there are a few considerations to take into account

- Schedule a downtime before taking a backup, this would help with data consistency for stateful applications
- For stateless applications built on the 12-factor principles and if application require to be highly available, consider having a global load balancer to route traffic to an active-active cluster and block traffic to the source cluster.
- Consider how pod to pod routing of services is designed, if a service mesh is being considered in the design as well
- For stateless applications built on the 12-factor principles check if a backup is necessary or if the application can be deployed to the target cluster with a pipeline in place
- Consider application retry mechanisms in place for events , message bus etc.

The steps give an overview of backing up all the resources in a cluster as well as backing up just a namespace in a cluster. This document does not go through the steps required to be considered before backing up a cluster.

The typical uses cases to perform a K8s resource backup is to facilitate application migrations from one cluster to another or between namespaces etc.

Step 1: Get kube config for the source cluster

```
pks login -a <pks api> -u <user> -p <password> -k  
pks get-credentials <cluster>
```

Alternatively

```
pks get-kubeconfig <cluster> -a <pks api> -u <user> -p <password> -k
```

E.g.

```
pks login -a pks.corp.local -u riaz -p VMware1! -k  
pks get-credentials ci-cluster  
pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k
```

```
ubuntu@cli-vm:~/velero$ pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k  
Fetching kubeconfig for cluster ci-cluster and user riaz.  
You can now use the kubeconfig for user riaz:  
$kubectl config use-context ci-cluster
```

Step 2: Set kubectl context to the source cluster

```
kubectl config use-context <source-cluster>
```

E.g.

```
kubectl config use-context ci-cluster
```

Step 3: Check all resources running on the source cluster

```
kubectl get ns
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get ns  
NAME      STATUS  AGE  
default   Active  32d  
kube-node-lease  Active  32d  
kube-public  Active  32d  
kube-system  Active  32d  
pks-system  Active  32d  
velero     Active  6d22h  
x1         Active  20d  
y1         Active  20d  
yelb       Active  3d19h  
z1         Active  20d
```

NOTE: apart from the default and system namespaces, yelb, x1, y1 and z1 exist

```
kubectl get po --all-namespaces
```

TKG Backup and Restore Clusters

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	busybox-7b87695f88-jgc4f	0/1	CrashLoopBackOff	6010	21d
kube-system	coredns-6fbcd956-c4n	1/1	Running	1	3d2d
kube-system	coredns-6fbcd956-kmcp	1/1	Running	1	3d2d
kube-system	kubernetes-dashboard-f5fc4ccc79f-csmb5	1/1	Running	1	3d2d
kube-system	metrics-server-zf85c59675-sjczr	0/1	Completed	0	32d
pks-system	cert-generator-11b35c51b71ea3086396a780dbf20b5cd695b25d-dbpxp7	2/2	Running	0	21d
pks-system	event-controller-7b96987577-kjczf	2/2	Running	0	21d
pks-system	fluent-bit-hxh69	2/2	Running	0	21d
pks-system	fluent-bit-wqgxw	2/2	Running	0	21d
pks-system	fluent-bit-wxpx	2/2	Running	0	21d
pks-system	metric-controller-66bb66498-69zsh	1/1	Running	0	21d
pks-system	observability-manager-7dd6c4cd-gg4q9	1/1	Running	1	32d
pks-system	sink-controller-5d76d8d546-sbghh	1/1	Running	0	21d
pks-system	telegraf-9cmrn	1/1	Running	0	21d
pks-system	telegraf-gnj73	1/1	Running	0	21d
pks-system	telegraf-zmvnx	1/1	Running	0	21d
pks-system	telegraf-zvqnt-58797bf64d-7skz5	2/2	Running	2	32d
pks-system	validator-7cb99cc-pzhrl	1/1	Running	0	21d
pks-system	vrops-cadvisor-xkj4p	1/1	Running	1	32d
pks-system	vrops-cadvisor-r43cp	1/1	Running	1	32d
pks-system	vrops-cadvisor-wfw4h	1/1	Running	1	32d
velero	restic-hmzfl	1/1	Running	1	32d
velero	restic-jff8c	1/1	Running	0	6d22h
velero	restic-s9fix	1/1	Running	0	6d22h
velero	velero-45d4c59-2c9rv	1/1	Running	0	6d22h
x1	service-a-84965f57cc-uhbrx	2/2	Running	187	20d
x1	service-a-84965f57cc-q72kh	2/2	Running	187	20d
y1	service-b-8eddc888c7-6r9kp	2/2	Running	187	20d
y1	service-b-8eddc888c7-jpwf9	2/2	Running	187	20d
yelb	redis-server-d	1/1	Running	0	3d19h
yelb	yelb-appserver-696b9668c4-7mrhs	1/1	Running	0	3d19h
yelb	yelb-db-0	1/1	Running	0	3d19h
yelb	yelb-db-1-66b575695-21qsw	1/1	Running	0	3d19h
z1	service-c-5cfc5c857-d9eqm	2/2	Running	187	20d
z1	service-c-5cfc5c857-jxqbh	2/2	Running	187	20d
z1	service-d-69ff59f4cb9-f94rs	2/2	Running	187	20d
z1	service-d-69ff59f4cb9-pkplt	2/2	Running	187	20d

NOTE: Note the pods running in the yelb , x1, y1 and z1 namespaces

kubectl get pv --all-namespaces

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
pvc-443e4fd9-a024-485d-b0fb-2273b668f908	2Gi	RWO	Delete	Bound	yelb/redis-pv-claim	thin-disk		3d19h
pvc-ed1b4b11-661a-450f-9d11-84c0acbbf1a0	5Gi	RWO	Delete	Bound	yelb/db-pv-claim	thin-disk		3d19h

NOTE: There are a couple of PVs in the yelb namespace (redis and yelb database)

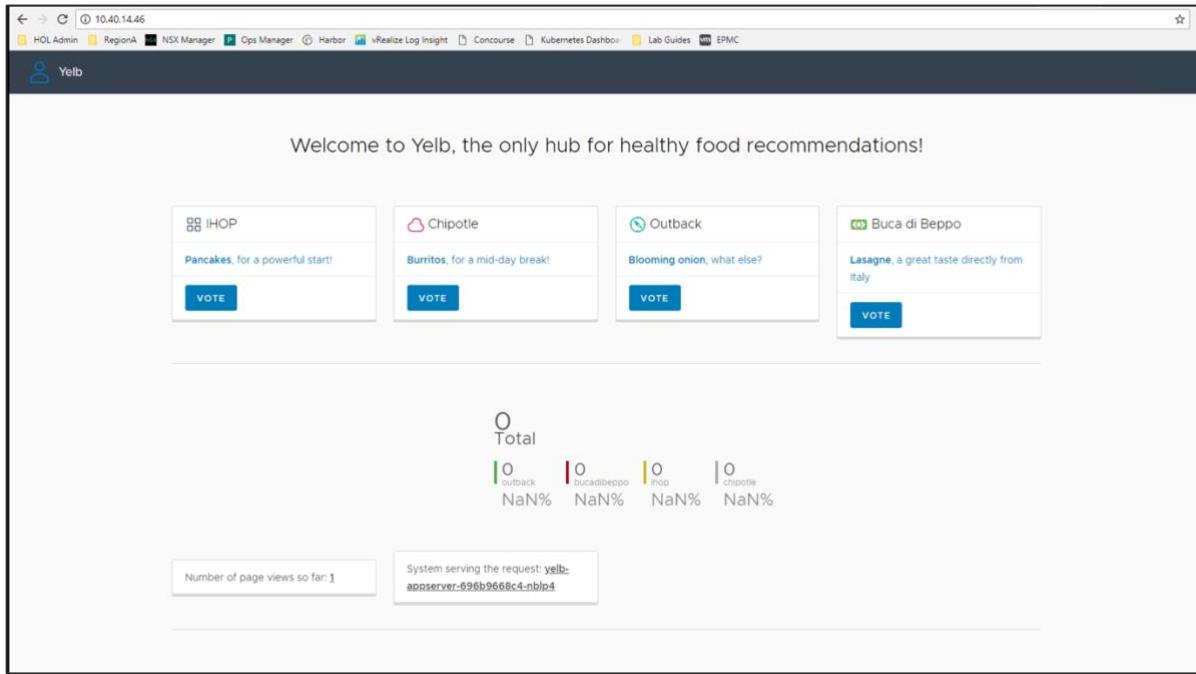
Step 4: Login to the Yelb application and make sure that the application is reachable

kubectl get svc -n yelb

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
redis-server	ClusterIP	10.100.200.113	<none>	6379/TCP	2m12s
yelb-appserver	ClusterIP	10.100.200.155	<none>	4567/TCP	2m11s
yelb-db	ClusterIP	10.100.200.158	<none>	5432/TCP	2m12s
yelb-ui	LoadBalancer	10.100.200.3	10.40.14.46	80:30561/TCP	2m11s

TKG Backup and Restore Clusters

Point the browser to the external-ip of yelb-ui service



Step 5: The Yelb application runs a yelb-db and the redis-server pods which are both stateful. All stateful pods need to be annotated.

NOTE: when using the velero vsphere plugin annotation is not required (Skip Step 5, Step 6 and Step 7)

Run the following to annotate each pod that contains a volume to back up

```
kubectl -n YOUR_POD_NAMESPACE annotate pod/YOUR_POD_NAME backup.velero.io/backup-volumes=YOUR_VOLUME_NAME_1,YOUR_VOLUME_NAME_2,...
```

Step 6: To find the volumes for the stateful pod, identify the stateful pod and describe it. For eg. in the Yelb deployment is the stateful pod

```
kubectl get po -n yelb
```

NAME	READY	STATUS	RESTARTS	AGE
redis-server-0	1/1	Running	0	3d19h
yelb-appserver-696b9668c4-7mrhs	1/1	Running	0	3d19h
yelb-db-0	1/1	Running	0	3d19h
yelb-ui-6665575695-2lqzw	1/1	Running	0	3d19h

The yelb application has two stateful pods (yelb-db and the redis-server)

```
kubectl describe po yelb-db-0 -n yelb
```

The volumes associated with this pod is

Volumes:

mysql-persistent-storage:

Type: PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)

ClaimName: db-pv-claim

ReadOnly: false

TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl describe po yelb-db-0 -n yelb
Name:           yelb-db-0
Namespace:      yelb
Priority:      0
PriorityClassName: <none>
Node:          Seb7adae-e071-4285-afdb-f9ab465db019/172.16.0.5
Start Time:    Thu, 04 Jun 2020 20:31:12 +0000
Labels:        app=yelb-db
               controller-revision-hash=yelb-db-5b6476565d
               statefulset.kubernetes.io/pod-name=yelb-db-0
               tier=back end db
Annotations:   backup.velero.io/backup-volumes: mysql-persistent-storage
Status:        Running
IP:            172.15.28.4
Controlled By: StatefulSet/yelb-db
Containers:
  yelb-db:
    Container ID: docker://9a9fc40a6aa68c50ce7df1966e8e8946d4c71e80a86bfff03c8988fa64d7fc4bf
    Image:         mrefere/yelb-db:0.3
    Image ID:     docker-pullable://mrefere/yelb-db@sha256:ddccf0943fe4f1146aac5c075ca8d263c5bc9c5b0f5c59cdfda2326f4813152a
    Port:          5432/TCP
    Host Port:    0/TCP
    State:        Running
    Started:     Thu, 04 Jun 2020 20:31:37 +0000
    Ready:        True
    Restart Count: 0
    Environment:
      PGDATA: /var/lib/postgresql/data/pgdata
    Mounts:
      /var/lib/postgresql/data from mysql-persistent-storage (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-9x96z (ro)
Conditions:
  Type        Status
  Initialized  True
  Ready       True
  ContainersReady  True
  PodScheduled  True
Volumes:
  mysql-persistent-storage:
    Type:      PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
    ClaimName: db-pv-claim
    ReadOnly:  false
    default-token-9x96z:
      Type:      Secret (a volume populated by a Secret)
      SecretName: default-token-9x96z
      Optional:   false
QoS Class:  BestEffort
Node-Selectors: <none>
Tolerations:  node.kubernetes.io/not-ready:NoExecute for 300s
               node.kubernetes.io/unreachable:NoExecute for 300s
Events:      <none>
```

kubectl describe po redis-server-0 -n yelb

The volumes associated with this pod is

Volumes:

redis-persistent-storage:

Type: PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
ClaimName: redis-pv-claim
ReadOnly: false

Step 7: Annotate the yelb-db-0 and redis-server-0 pods

kubectl -n yelb annotate pod/< yelb db name for eg. yelb-db-0 > backup.velero.io/backup-volumes=<volume name>

e.g.

kubectl -n yelb annotate pod/yelb-db-0 backup.velero.io/backup-volumes=mysql-persistent-storage

kubectl -n yelb annotate pod/redis-server-0 backup.velero.io/backup-volumes=redis-persistent-storage

Step 8: Login to the Yelb application and vote for some of the restaurants listed. The voting data is persisted in the Yelb database, this data will validate the state of the application when restored. Keep a tab on the votes.

The screenshot shows the Yelb application running on a web browser. The URL in the address bar is 10.40.14.44. The page title is 'Yelb'. The main content area displays four restaurant cards: IHOP, Chipotle, Outback, and Buca di Beppo. Each card includes a description and a 'VOTE' button. Below the cards is a pie chart titled '15 Total' showing the distribution of votes across four categories: outback (27%), bucadibeppo (20%), hop (33%), and chipotle (20%). At the bottom left, there is a box stating 'Number of page views so far: 2'. At the bottom right, there is a box stating 'System serving the request: yelb-appserver-696b9668c4-4fvn9'.

Candidate	Count	Percentage
outback	4	27%
bucadibeppo	3	20%
hop	5	33%
chipotle	3	20%

Cluster Backup Using the Restic Plugin

Step 9: Create a backup all the resources in a cluster

```
cd ~/velero/velero-v1.4.0-linux-amd64
```

```
./velero create backup <BACKUP NAME>
```

E.g.

```
./velero create backup sourceclusterbk
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero create backup sourceclusterbk
Backup request "sourceclusterbk" submitted successfully.
Run `velero backup describe sourceclusterbk` or `velero backup logs sourceclusterbk` for more details.
```

Step 10: Check status of the backup. The output describes the status of the backup. We have taken a complete backup and hence all resources and PVs are backed up.

```
./velero backup describe <BACKUP NAME>
```

E.g.

```
./velero backup describe sourceclusterbk
```

```
Ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero backup describe sourceclusterbk
Name: sourceclusterbk
Namespace: velero
Labels: velero.io/storage-location=default
Annotations: velero.io/source-cluster-k8s-gitversion=v1.15.5
velero.io/source-cluster-k8s-major-version=1
velero.io/source-cluster-k8s-minor-version=15
Phase: Completed
Namespaces:
Included: *
Excluded: <none>
Resources:
Included: *
Excluded: <none>
Cluster-scoped: auto
Label selector: <none>
Storage Location: default
Velero-Native Snapshot PVs: auto
TTL: 720h0m0s
Hooks: <none>
Backup Format Version: 1
Started: 2020-06-09 04:20:33 +0000 UTC
Completed: 2020-06-09 04:20:49 +0000 UTC
Expiration: 2020-07-09 04:20:33 +0000 UTC
Total items to be backed up: 570
Items backed up: 570
Velero-Native Snapshots: <none included>
Restic Backups (specify --details for more information):
Completed: 2
```

Note: The restic backups – these are backups of the pv's in this case the yelb-db and redis server

Step 11: Login to Minio and check if the backup and restic folders have been created. Use <http://<minio-service-external-ip>>, e.g. <http://10.40.14.136>.

To find the ip of minio check the IP under the “External-IP” section, point your browser to the location <external-ip>. The Minio application should be accessible

```
kubectl get svc -n minio
```

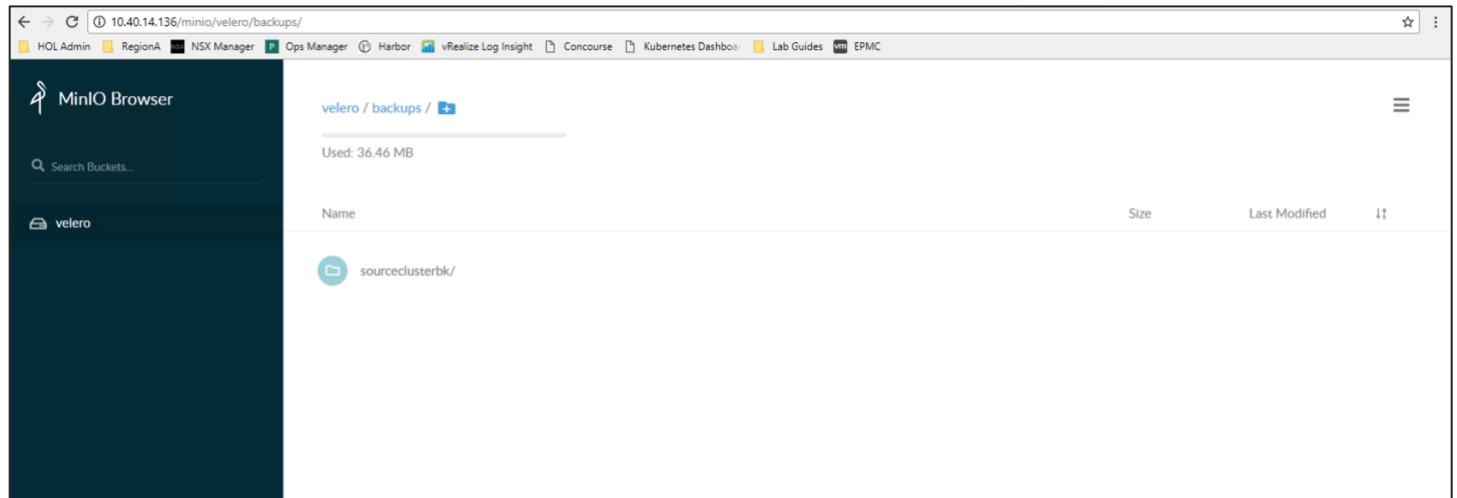
TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero$ kubectl get svc -n minio
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE
minio-frontend-lb   LoadBalancer   10.100.200.3   10.40.14.136   80:30568/TCP   18s
minio-release     ClusterIP    10.100.200.107  <none>         9000/TCP      13m
```

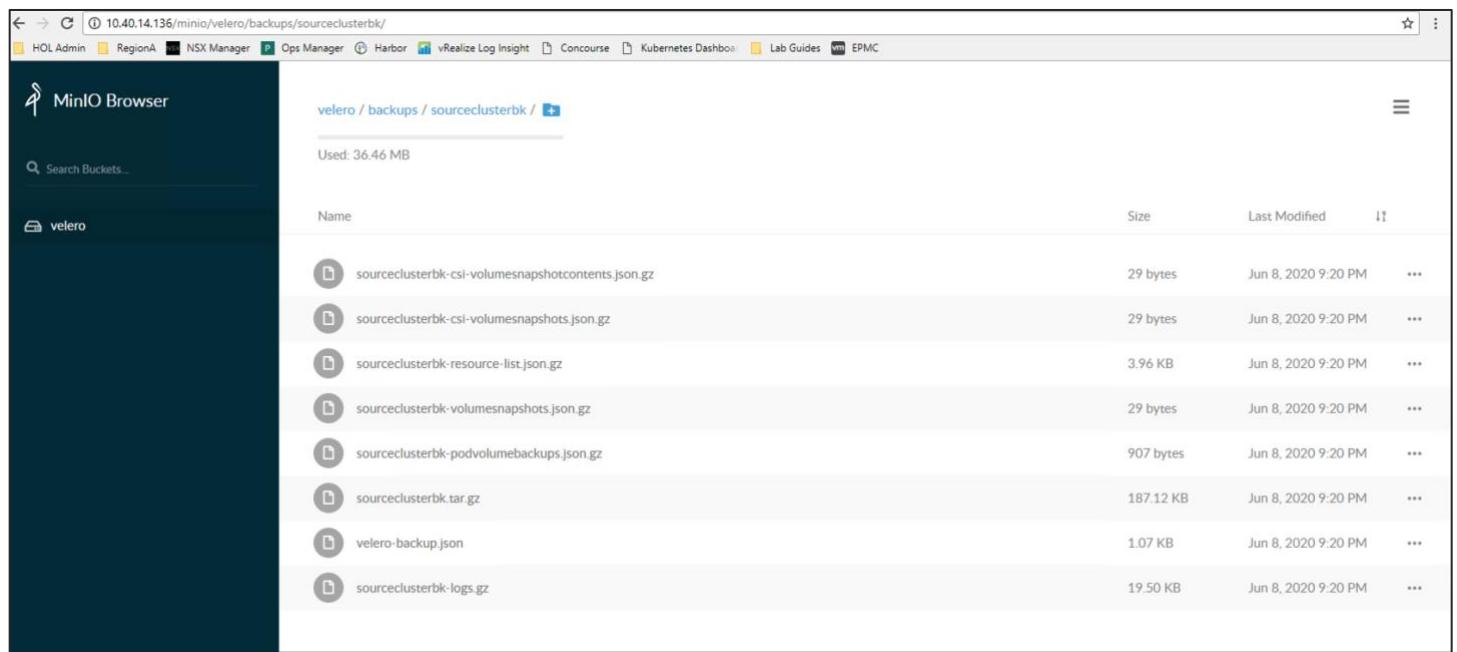
The screenshot shows a web-based interface for managing MinIO buckets. The left sidebar has a dark theme with a search bar for 'Search Buckets...' and a folder icon labeled 'velero'. The main area shows a list of buckets under the path 'velero /'. There are two entries: 'backups/' and 'restic/'. Below the list, it says 'Used: 36.46 MB'. The top navigation bar includes links for HOL Admin, RegionA, NSX Manager, Ops Manager, Harbor, vRealize Log Insight, Concourse, Kubernetes Dashboard, Lab Guides, and EPMC.

The backup folder contains the resource backup's and the restic contains the persistent volume backup's, the pv backups are referenced within the backup.

TKG Backup and Restore Clusters

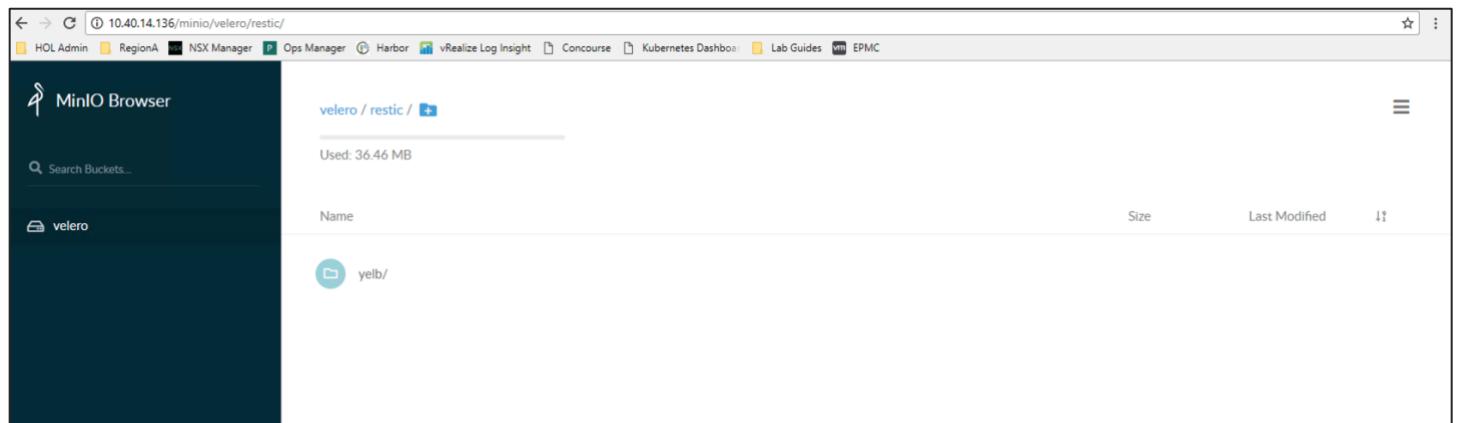


The screenshot shows the MinIO Browser interface with the URL `10.40.14.136/minio/velero/backups/`. The left sidebar has a dark theme with the MinIO logo and a search bar. The main area shows a single folder entry: `sourceclusterbk/`.



The screenshot shows the MinIO Browser interface with the URL `10.40.14.136/minio/velero/backups/sourceclusterbk/`. The left sidebar has a dark theme with the MinIO logo and a search bar. The main area lists several backup files and artifacts:

Name	Size	Last Modified	Actions
sourceclusterbk-csi-volumesnapshotcontents.json.gz	29 bytes	Jun 8, 2020 9:20 PM	...
sourceclusterbk-csi-volumesnapshots.json.gz	29 bytes	Jun 8, 2020 9:20 PM	...
sourceclusterbk-resource-list.json.gz	3.96 KB	Jun 8, 2020 9:20 PM	...
sourceclusterbk-volumesnapshots.json.gz	29 bytes	Jun 8, 2020 9:20 PM	...
sourceclusterbk-podvolumebackups.json.gz	907 bytes	Jun 8, 2020 9:20 PM	...
sourceclusterbk.tar.gz	187.12 KB	Jun 8, 2020 9:20 PM	...
velero-backup.json	1.07 KB	Jun 8, 2020 9:20 PM	...
sourceclusterbk-logs.gz	19.50 KB	Jun 8, 2020 9:20 PM	...



The screenshot shows the MinIO Browser interface with the URL `10.40.14.136/minio/velero/restic/`. The left sidebar has a dark theme with the MinIO logo and a search bar. The main area shows a single folder entry: `yelb/`.

TKG Backup and Restore Clusters

The screenshot shows the MinIO Browser interface. The URL in the address bar is 10.40.14.136/minio/velero/restic/yelb/. The browser's navigation bar includes links for HOL Admin, RegionA, NSX Manager, Ops Manager, Harbor, vRealize Log Insight, Concourse, Kubernetes Dashboard, Lab Guides, and EPMC. The main content area is titled "velero / restic / yelb /". It displays a list of objects with the following details:

Name	Size	Last Modified	...
data/			
Index/			
keys/			
snapshots/			
config	155 bytes	Jun 8, 2020 9:20 PM	...

Step 12: Login to the Yelp application and more votes to the restaurants listed. The voting data is persisted in the Yelp database, this data will validate the state of the application when restored.

The screenshot shows the Yelp application interface. The URL in the address bar is 10.40.14.46. The browser's navigation bar includes links for HOL Admin, RegionA, NSX Manager, Ops Manager, Harbor, vRealize Log Insight, Concourse, Kubernetes Dashboard, Lab Guides, and EPMC. The main content area is titled "Yelp". It displays a welcome message: "Welcome to Yelp, the only hub for healthy food recommendations!". Below this, there are four cards for different restaurants:

- IHOP**: Pancakes, for a powerful start! [VOTE](#)
- Chipotle**: Burritos, for a mid-day break! [VOTE](#)
- Outback**: Blooming onion, what else? [VOTE](#)
- Buca di Beppo**: Lasagne, a great taste directly from Italy [VOTE](#)

Below the cards is a pie chart titled "38 Total" showing the distribution of votes:

Restaurant	Count	Percentage
outback	12	32%
bucadibeppo	8	21%
ihop	14	37%
chipotle	4	11%

Cluster Backup Using the vSphere plugin

Step 1: Create a Volume snapshot location. This Volume Snapshot location is referenced when a backup is taken

```
cd ~/velero/velero-v1.4.0-linux-amd64
```

```
./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere
Snapshot volume location "vsphere-snap-loc" configured successfully.
```

Step 2: Create a cluster backup

```
./velero backup create srccluster-snap-backup --snapshot-volumes --volume-snapshot-locations
vsphere-snap-loc
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero backup create srccluster-snap-backup --snapshot-volumes --volume-snapshot-locations vsphere-snap-loc
Backup request "srccluster-snap-backup" submitted successfully.
Run `velero backup describe srccluster-snap-backup` or `velero backup logs srccluster-snap-backup` for more details.
```

Step 3: Check status of the backup. The output describes the status of the backup. We have taken a complete backup and hence all resources and PVs are backed up.

```
./velero backup describe <BACKUP NAME>
```

For more details on the backup

```
./velero backup describe <BACKUP NAME> --details
```

E.g.

```
./velero backup describe srccluster-snap-backup
```

TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero backup describe srccluster-snap-backup
Name:           srccluster-snap-backup
Namespace:      velero
Labels:         velero.io/storage-location=default
Annotations:   velero.io/source-cluster-k8s-gitversion=v1.16.7+vmware.1
               velero.io/source-cluster-k8s-major-version=1
               velero.io/source-cluster-k8s-minor-version=16

Phase:          Completed

Namespaces:
  Included:    *
  Excluded:   <none>

Resources:
  Included:    *
  Excluded:   <none>
  Cluster-scoped: auto

Label selector: <none>

Storage Location: default

Velero-Native Snapshot PVs: true

TTL: 720h0m0s

Hooks: <none>

Backup Format Version: 1

Started: 2020-07-08 17:00:11 +0000 UTC
Completed: 2020-07-08 17:00:46 +0000 UTC

Expiration: 2020-08-07 17:00:11 +0000 UTC

Total items to be backed up: 575
Items backed up: 575

Velero-Native Snapshots: 2 of 2 snapshots completed successfully (specify --details for more information)
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$
```

NOTE: Event if the velero backup status shows complete the backup would not have been complete. When the vsphere plugin is used there is a lot of operations that happen in the background. During the backup there are a number of operations that occur in Vsphere.

Step 4: Monitor the uploads

```
kubectl -n velero get uploads
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl -n velero get uploads
NAME          AGE
upload-2d9203ed-ac81-4e98-b468-15ca21658766  10m
upload-2dcc8590-d02b-4bf2-9222-f51562c971c5  19m
upload-75423580-9dcc-405d-8e34-7b324a226022  19m
upload-864510ed-73f1-4af8-ba23-5be409eb6b50  19m
upload-8671eb8b-f703-4525-8797-81d3e0e0fd9f  19m
upload-c6d903d2-ed4a7-41b2-b6de-6f2b339bf009  10m
upload-df77bd51-56db-45e7-9ae6-225e8fb6a3f7  10m
upload-fd1460d6-710a-4c0b-9b1b-422fa5373078  10m
```

NOTE: Until the uploads are complete do not perform a restore operation

Step 5: Login to Minio and check if the backup and plugins folders have been created.

To find the ip of minio check the IP under the “External-IP” section, point your browser to the location <external-ip>. The Minio application should be accessible

```
kubectl get svc -n minio
```

ubuntu@cli-vm:~/velero\$ kubectl get svc -n minio					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
minio-frontend-lb	LoadBalancer	10.100.200.3	10.40.14.136	80:30568/TCP	18s
minio-release	ClusterIP	10.100.200.107	<none>	9000/TCP	13m

The screenshot shows the MinIO Browser interface. At the top, there is a navigation bar with various links: HOL Admin, RegionA, NSX Manager, Ops Manager, Harbor, vRealize Log Insight, Concourse, Kubernetes Dashboard, Lab Guides, and EPMC. Below the navigation bar, the title "MinIO Browser" is displayed, along with a search bar labeled "Search Buckets..." and a "velero" folder icon.

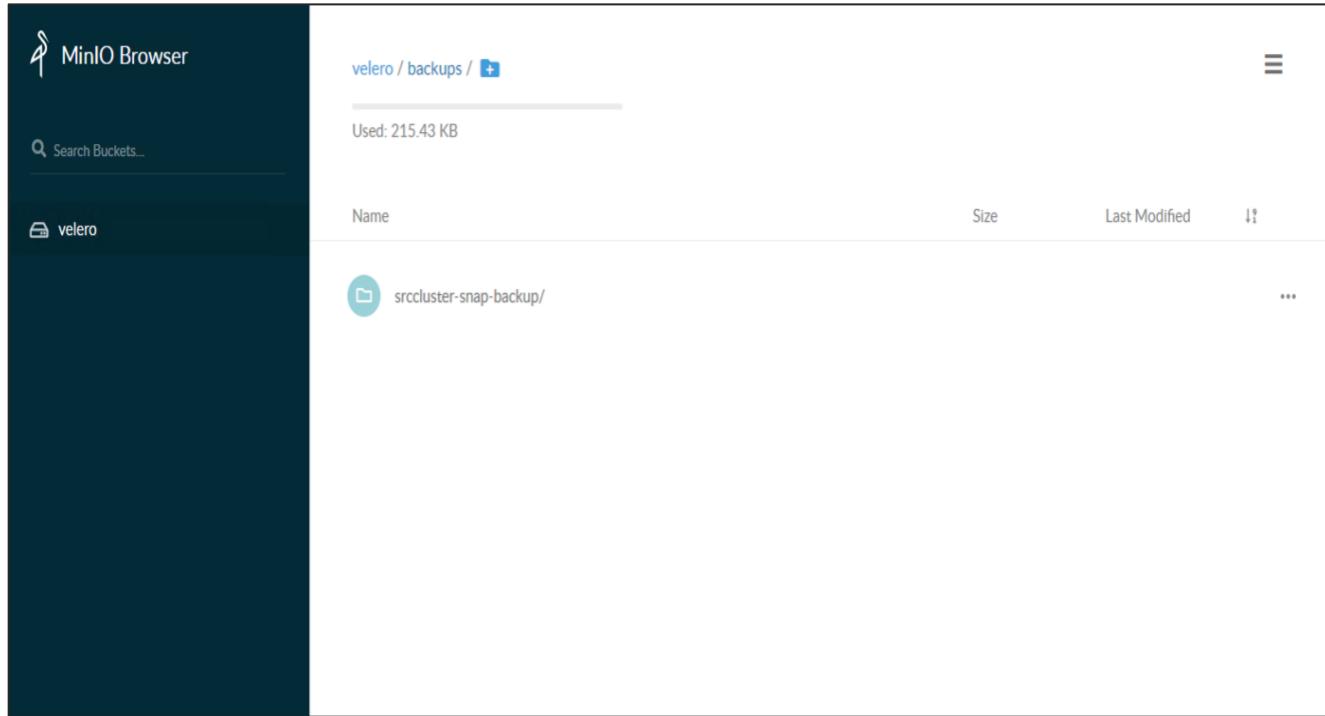
The main area shows the contents of the "velero" bucket. It displays a summary: "Used: 215.43 KB". Below this, a table lists the contents of the bucket:

Name	Size	Last Modified
backups/		...
plugins/		...

TKG Backup and Restore Clusters

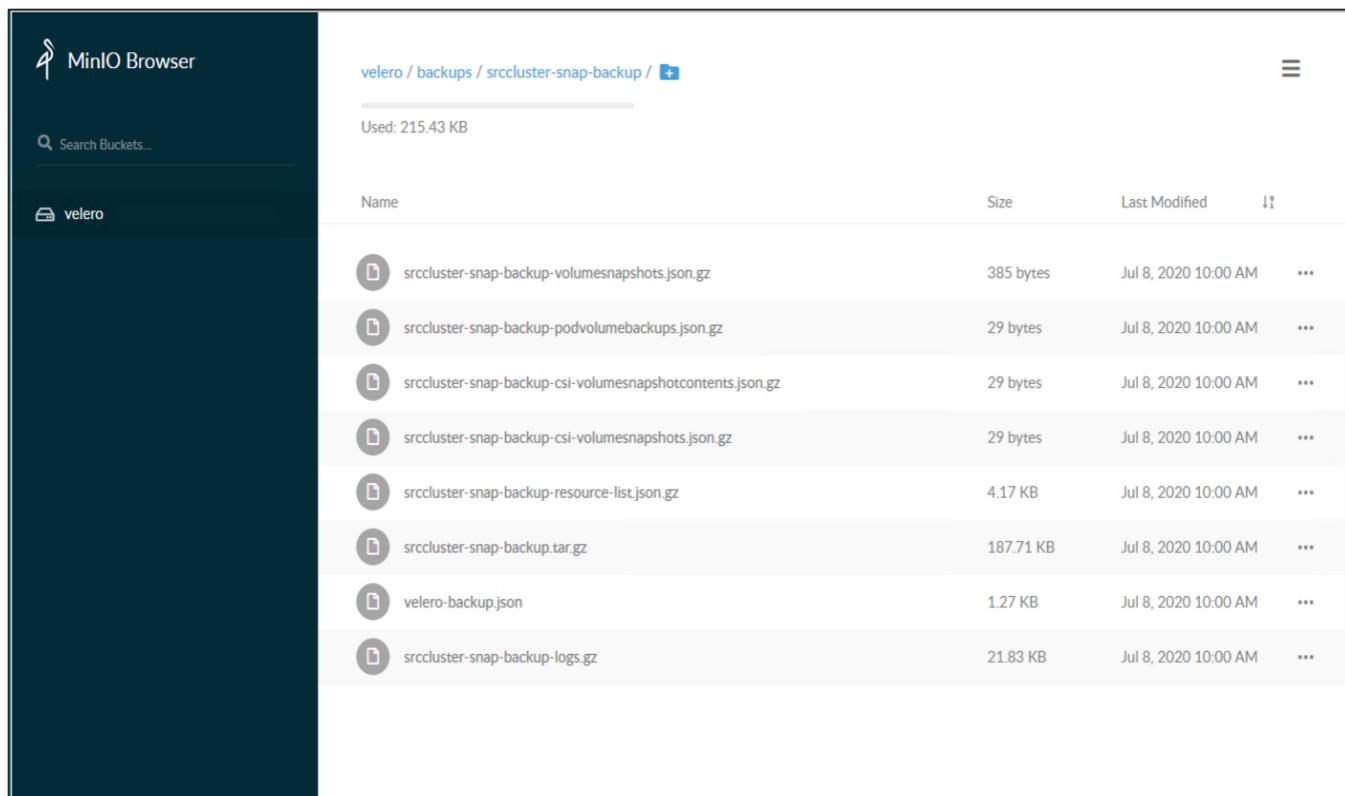
The screenshot shows the MinIO Browser interface. On the left, there is a sidebar with a search bar labeled "Search Buckets..." and a folder icon labeled "velero". The main area displays a file tree under the path "velero / plugins / vsphere-astrolabe-repo / ivd /". The tree shows three sub-folders: "data/", "md/", and "peinfo/". The total usage is listed as "Used: 215.43 KB". The interface includes a header with the current path and a "+" button, and a menu icon in the top right corner.

TKG Backup and Restore Clusters



The screenshot shows the MinIO Browser interface. The left sidebar has a dark theme with the title "MinIO Browser" and a search bar labeled "Search Buckets...". Under the "velero" bucket, there is a folder named "srccluster-snap-backup/". The main panel displays the contents of this folder.

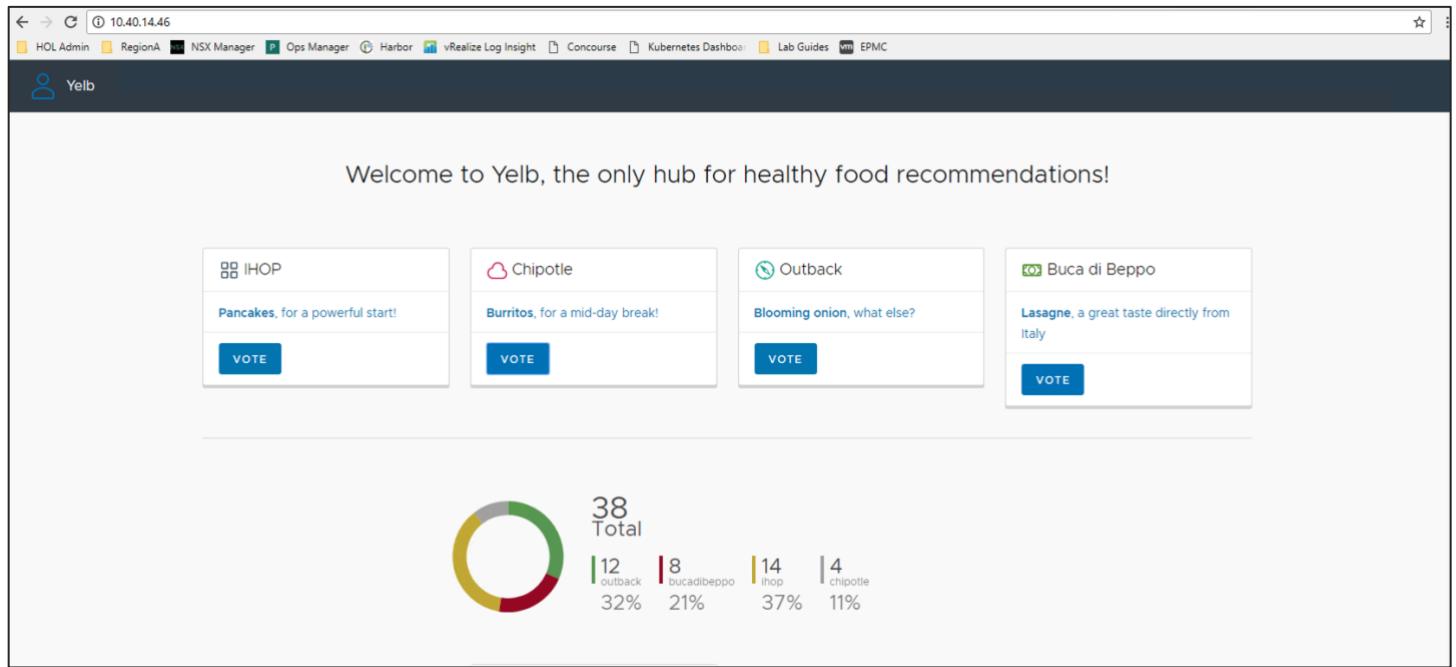
Name	Size	Last Modified	Actions
srccluster-snap-backup/			...



The screenshot shows the MinIO Browser interface, similar to the first one but at a deeper level of the backup hierarchy. The left sidebar is identical. The main panel shows the contents of the "srccluster-snap-backup/" folder.

Name	Size	Last Modified	Actions
srccluster-snap-backup-volumesnapshots.json.gz	385 bytes	Jul 8, 2020 10:00 AM	...
srccluster-snap-backup-podvolumebackups.json.gz	29 bytes	Jul 8, 2020 10:00 AM	...
srccluster-snap-backup-csi-volumesnapshotcontents.json.gz	29 bytes	Jul 8, 2020 10:00 AM	...
srccluster-snap-backup-csi-volumesnapshots.json.gz	29 bytes	Jul 8, 2020 10:00 AM	...
srccluster-snap-backup-resource-list.json.gz	4.17 KB	Jul 8, 2020 10:00 AM	...
srccluster-snap-backup.tar.gz	187.71 KB	Jul 8, 2020 10:00 AM	...
velero-backup.json	1.27 KB	Jul 8, 2020 10:00 AM	...
srccluster-snap-backup-logs.gz	21.83 KB	Jul 8, 2020 10:00 AM	...

Step 6: Login to the Yelb application and more votes to the restaurants listed. The voting data is persisted in the Yelb database, this data will validate the state of the application when restored.



Namespace Backup Using the Restic Plugin

Step 1: Create a backup of the yelb namespace

```
cd ~/velero/velero-v1.4.0-linux-amd64  
./velero backup create <BACKUP NAME> --include-namespaces <NAMESPACE1>
```

E.g.

```
./velero backup create yelbbkup --include-namespaces yelb
```

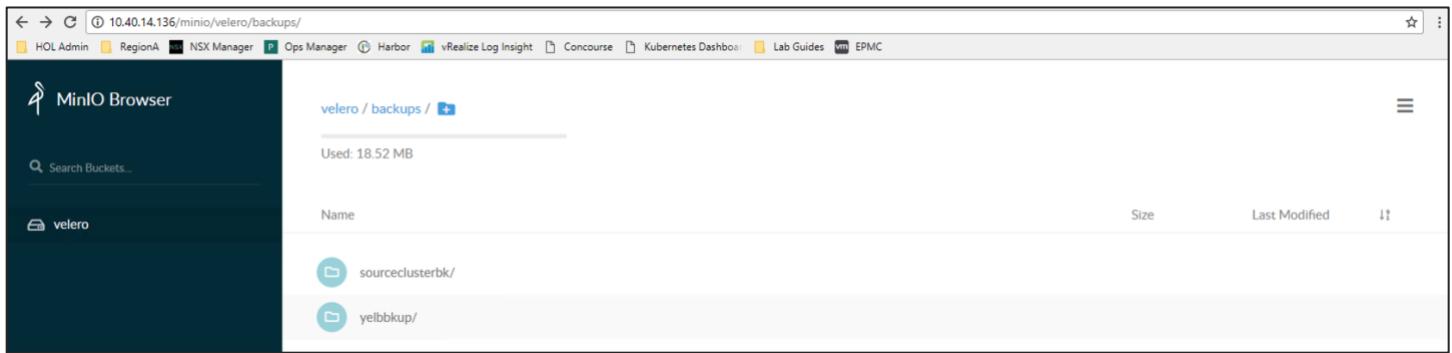
Step 2: Check status of the backup

```
./velero backup describe yelbbkup
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero backup describe yelbbkup  
Name: yelbbkup  
Namespace: velero  
Labels:  
Annotations: velero.io/storage-location=default  
velero.io/source-cluster-k8s-gitversion=v1.15.5  
velero.io/source-cluster-k8s-major-version=1  
velero.io/source-cluster-k8s-minor-version=15  
Phase: Completed  
Namespaces:  
Included: yelb  
Excluded: <none>  
Resources:  
Included: *  
Excluded: <none>  
Cluster-scoped: auto  
Label selector: <none>  
Storage Location: default  
Velero-Native Snapshot PVs: auto  
TTL: 720h0m0s  
Hooks: <none>  
Backup Format Version: 1  
Started: 2020-06-09 04:32:33 +0000 UTC  
Completed: 2020-06-09 04:32:39 +0000 UTC  
Expiration: 2020-07-09 04:32:33 +0000 UTC  
Total items to be backed up: 53  
Items backed up: 53  
Velero-Native Snapshots: <none included>  
Restic Backups (specify --details for more information):  
Completed: 2
```

The backup describes the status , the resources etc. In this case as shown , only the resources associated with namespace yelb is backed-up. The Restic backups describe the number of PV's that have been backed up.

Step 3: Login to minio and check if the backup has been created.



A screenshot of the MinIO Browser interface. The URL in the address bar is 10.40.14.136/minio/velero/backups/. The page displays a list of backups under the 'velero / backups /' path. The list shows two entries: 'sourceclusterbk/' and 'yelbbkup/'. Below the list, it says 'Used: 18.52 MB'. The interface includes a sidebar with a 'MinIO Browser' logo and a search bar labeled 'Search Buckets...'. The top navigation bar contains various links like HOL Admin, RegionA, NSX Manager, Ops Manager, Harbor, vRealize Log Insight, Concourse, Kubernetes Dashboard, Lab Guides, and EPMC.

Step 4: Other options for backup

```
./velero backup create planes --selector app=yelb
```

Check Velero documentation <https://velero.io/docs/v1.4/> for other options

Namespace Backup Using the vSphere plugin

Step 1: Create a Volume snapshot location. This Volume Snapshot location is referenced when a backup is taken

NOTE: If a snapshot location has already been created this step is optional

```
cd ~/velero/velero-v1.4.0-linux-amd64  
. ./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere  
Snapshot volume location "vsphere-snap-loc" configured successfully.
```

Step 2: Create a namespace backup

```
./velero backup create yelb-vspheresnap-bkp --include-namespaces=yelb --snapshot-volumes --volume-snapshot-locations vsphere-snap-loc
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero backup create yelb-vspheresnap-bkp --include-namespaces=yelb --snapshot-volumes --volume-snapshot-locations vsphere-snap-loc  
Backup request "yelb-vspheresnap-bkp" submitted successfully.  
Run `velero backup describe yelb-vspheresnap-bkp` or `velero backup logs yelb-vspheresnap-bkp` for more details.
```

Step 3: Check status of the backup

```
./velero backup describe yelb-vspheresnap-bkp --details
```

```
- yelb/redis-server-0.161fddc04e981821
- yelb/redis-server-0.161fddc0dea8378d
- yelb/redis-server-0.161fddc3d9d76c06
- yelb/redis-server-0.161fddc3ea9defee
- yelb/redis-server-0.161fddc4019fe300
- yelb/redis-server.161fddc04ceb30f4
- yelb/yelb-appserver-794d7c9458-cblfn.161fddc0573dbf31
- yelb/yelb-appserver-794d7c9458-cblfn.161fddc113f84b07
- yelb/yelb-appserver-794d7c9458-cblfn.161fddc11ea7ceb0
- yelb/yelb-appserver-794d7c9458-cblfn.161fddc135a79906
- yelb/yelb-appserver-794d7c9458.161fddc056207e71
- yelb/yelb-appserver.161fddc05516668a
- yelb/yelb-db-0.161fddc0526e1081
- yelb/yelb-db-0.161fddc23e4a39bf
- yelb/yelb-db-0.161fddc59e98da85
- yelb/yelb-db-0.161fddcc1ccb2087
- yelb/yelb-db-0.161fddce111d5a04
- yelb/yelb-db-0.161fddce229e1360
- yelb/yelb-db.161fddc0506e7083d
- yelb/yelb-ui-79c68df689-r72hr.161fddc04968bce6
- yelb/yelb-ui-79c68df689-r72hr.161fddc0f90e255b
- yelb/yelb-ui-79c68df689-r72hr.161fddc6c35d8f95
- yelb/yelb-ui-79c68df689-r72hr.161fddc9687bd145
- yelb/yelb-ui-79c68df689-r72hr.161fddc98668e3be
- yelb/yelb-ui-79c68df689.161fddc048232b77
- yelb/yelb-ui.161fddc0470d1430

v1/Namespace:
- yelb

v1/PersistentVolume:
- pvc-9c5dc2c9-084d-4485-bc4f-a327ca47bd9a
- pvc-e6166d4e-a21c-44d6-89f2-c3dec5d92199

v1/PersistentVolumeClaim:
- yelb/db-pv-claim
- yelb/redis-pv-claim

v1/Pod:
- yelb/redis-server-0
- yelb/yelb-appserver-794d7c9458-cblfn
- yelb/yelb-db-0
- yelb/yelb-ui-79c68df689-r72hr

v1/Secret:
- yelb/default-token-p8jj9

v1/Service:
- yelb/redis-server
- yelb/yelb-appserver
- yelb/yelb-db
- yelb/yelb-ui

v1/ServiceAccount:
- yelb/default

Velero-Native Snapshots:
pvc-e6166d4e-a21c-44d6-89f2-c3dec5d92199:
  Snapshot ID:      ivd:2f2dc9a4-aee2-47a4-a909-5db8b4f837b9:26662c39-642a-4179-af29-a110e1d0292c
  Type:             ivd
  Availability Zone:
  IOPS:             100
pvc-9c5dc2c9-084d-4485-bc4f-a327ca47bd9a:
  Snapshot ID:      ivd:15091bc2-0844-4428-b3d5-69399c814620:23500ea7-e487-4313-91d7-27a6f7297c4c
  Type:             ivd
  Availability Zone:
  IOPS:             100
```

NOTE: Even if the velero backup status shows ‘complete’, the backup would not have been complete. When the vSphere plugin is used there is a lot of operations that happen in the background (vSphere etc.).

Step 4: Monitor the uploads

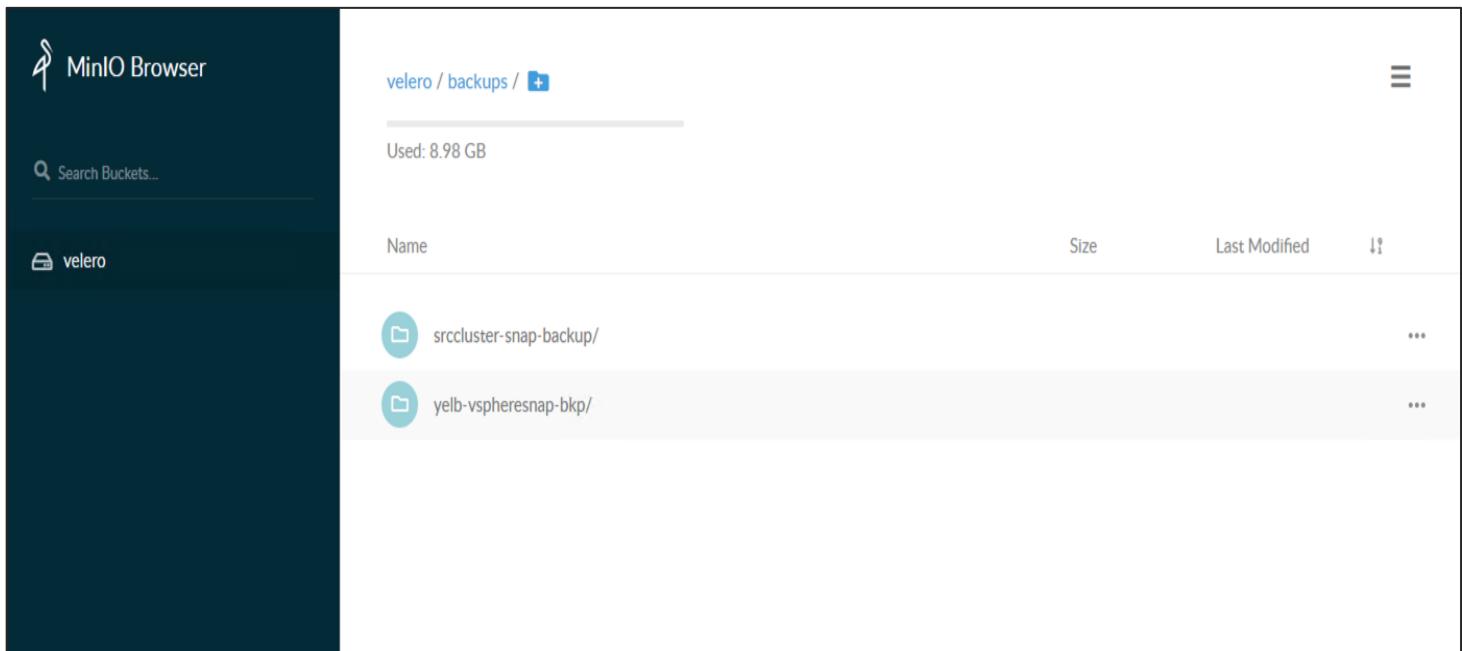
```
kubectl -n velero get uploads
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl -n velero get uploads
NAME                                AGE
upload-23500ea7-e487-4313-91d7-27a6f7297c4c  2m34s
upload-26662c39-642a-4179-af29-a110e1d0292c  2m45s
upload-2d9203ed-ac81-4e98-b468-15ca21658766  17m
upload-2dcc8590-d02b-4bf2-9222-f51562c971c5  26m
upload-75423580-9dcc-405d-8e34-7b324a226022  26m
upload-864510ed-73f1-4afd-ba23-5be409eb6b50  26m
upload-8671ebef-f703-4525-8797-81d3e0e0fd9f  26m
upload-c6d903d2-edaa7-41b2-b6de-6f2b339bf009  17m
upload-df77bd51-56db-45e7-9ae6-225e8fb6a3f7  17m
upload-fd1460d6-710a-4c0b-9b1b-422fa5373078  17m
```

NOTE: Until the uploads are complete do not perform a restore operation

Step 5: Check vCenter to make sure that none of the operations related to the backup are running

Step 6: Login to Minio and check if the backup exists



Name	Size	Last Modified
srccluster-snap-backup/		...
yelb-vspheresnap-bkp/		...

Scheduled Backups

Schedule Backup Using the Restic Plugin

The schedule operation allows you to back up your data at recurring intervals. The first backup is performed when the schedule is first created, and subsequent backups happen at the schedule's specified interval. These intervals are specified by a Cron expression.

Scheduled backups are point in time backups and can be used for disaster recovery use cases.

Step 1: Login to the Yelb application and more votes to the restaurants listed. The voting data is persisted in the Yelb database it will validate the state of the application when restored.

The screenshot shows the Yelb application running on a host with IP 10.40.14.46. The interface includes a navigation bar with links like HOL Admin, RegionA, NSX Manager, Ops Manager, vRealize Log Insight, Concourse, Kubernetes Dashboard, Lab Guides, and EPMC. The main content area displays four restaurant cards: IHOP, Chipotle, Outback, and Buca di Beppo, each with a 'VOTE' button. Below the cards is a pie chart titled '47 Total' showing the distribution of votes: Outback (36%), Buca di Beppo (17%), IHOP (38%), and Chipotle (8.51%). At the bottom, there are two informational boxes: 'Number of page views so far: 4' and 'System serving the request: yelb-appserver-696b9668c4-nblp4'.

Step 2: Create a backup scheduler

```
./velero schedule create <scheduler-name> --schedule <cron-expression>
```

E.g.

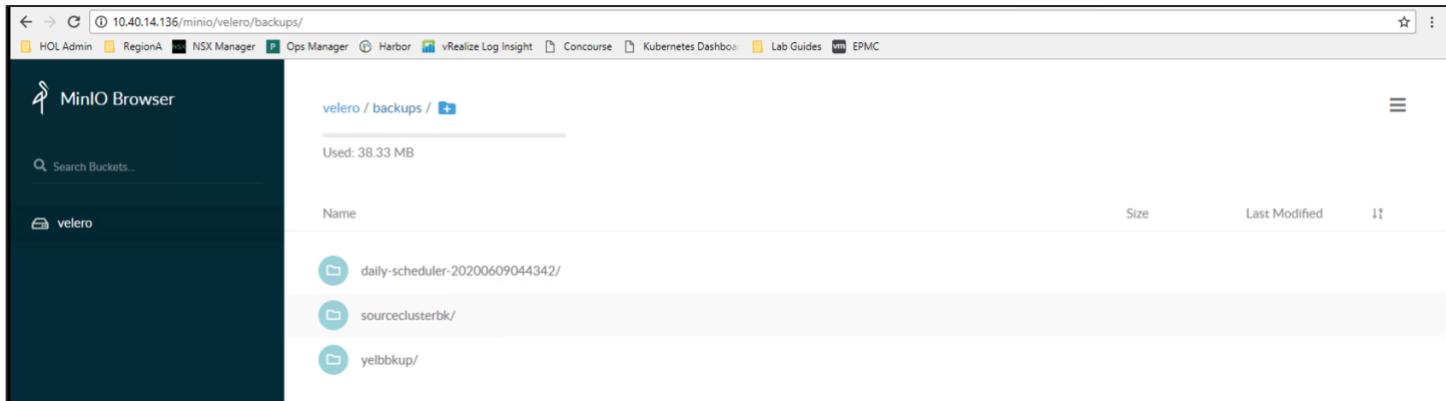
To create a daily scheduler (use <https://crontab.guru/every-15-minutes> for cron expressions if needed)

```
./velero schedule create daily-scheduler --schedule="@every 24h" --ttl 24h0m0s
```

Note: The TTL flag allows the user to specify the backup retention period with the value specified in hours, minutes, and seconds in the form --ttl 24h0m0s.

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero schedule create daily-scheduler --schedule="@every 24h" --ttl 24h0m0s
Schedule "daily-scheduler" created successfully.
```

Step 3: Login to Minio and check if the backup has been created. When creating a scheduler, the first backup will be created soon as the schedule is submitted.



Name	Size	Last Modified
daily-scheduler-20200609044342/		
sourceclusterbk/		
yelbbkup/		

Note: The state data of the application before each backup was performed was as below

Schedule Backup Using the vSphere plugin

Step 1: Login to the Yelb application and more votes to the restaurants listed. The voting data is persisted in the Yelb database , this data will validate the state of the application when restored.

The screenshot shows the Yelb application running on a web browser. The URL bar indicates the site is at 10.40.14.46. The top navigation bar includes links for HOL Admin, RegionA, NSX Manager, Ops Manager, Harbor, vRealize Log Insight, Concourse, Kubernetes Dashboard, Lab Guides, and EPMC. The main content area has a dark header with the text "Yelb". Below the header, a message reads "Welcome to Yelb, the only hub for healthy food recommendations!". There are four cards, each representing a restaurant with a logo, a name, a description, and a "VOTE" button:

- IHOP**: Pancakes, for a powerful start! (VOTE)
- Chipotle**: Burritos, for a mid-day break! (VOTE)
- Outback**: Blooming onion, what else? (VOTE)
- Buca di Beppo**: Lasagne, a great taste directly from Italy (VOTE)

Below the cards is a donut chart titled "Total" with the value 47. The chart shows the distribution of votes by restaurant:

Restaurant	Number of Votes	Percentage
outback	17	36%
bucadibepo	18	38%
ihop	18	38%
chipotle	4	8.51%

At the bottom left, a box displays "Number of page views so far: 4". At the bottom right, a box displays "System serving the request: yelb-appserver-696b9668c4-nbjp4".

Step 2: Create a Volume snapshot location. This Volume Snapshot location is referenced when a backup is taken

NOTE: If a snapshot location has already been created this step is optional

```
cd ~/velero/velero-v1.4.0-linux-amd64  
./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere  
Snapshot volume location "vsphere-snap-loc" configured successfully.
```

Step 3: Create a backup scheduler

```
./velero schedule create <scheduler-name> --schedule <cron-expression> --snapshot-volumes --  
volume-snapshot-locations vsphere-snap-loc
```

E.g.

To create a daily scheduler (use <https://crontab.guru/every-15-minutes> for cron expressions if needed)

```
./velero schedule create daily-scheduler --schedule="@every 24h" --ttl 24h0m0s --snapshot-volumes --  
volume-snapshot-locations vsphere-snap-loc
```

Note: The TTL flag allows the user to specify the backup retention period with the value specified in hours, minutes, and seconds in the form --ttl 24h0m0s.

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero schedule create daily-scheduler --schedule="@every 24h" --ttl 24h0m0s --snapshot-volumes --volume-snapshot-locations vsphere-snap-loc  
Schedule "daily-scheduler" created successfully.
```

Step 4: Monitor the uploads

```
kubectl -n velero get uploads
```

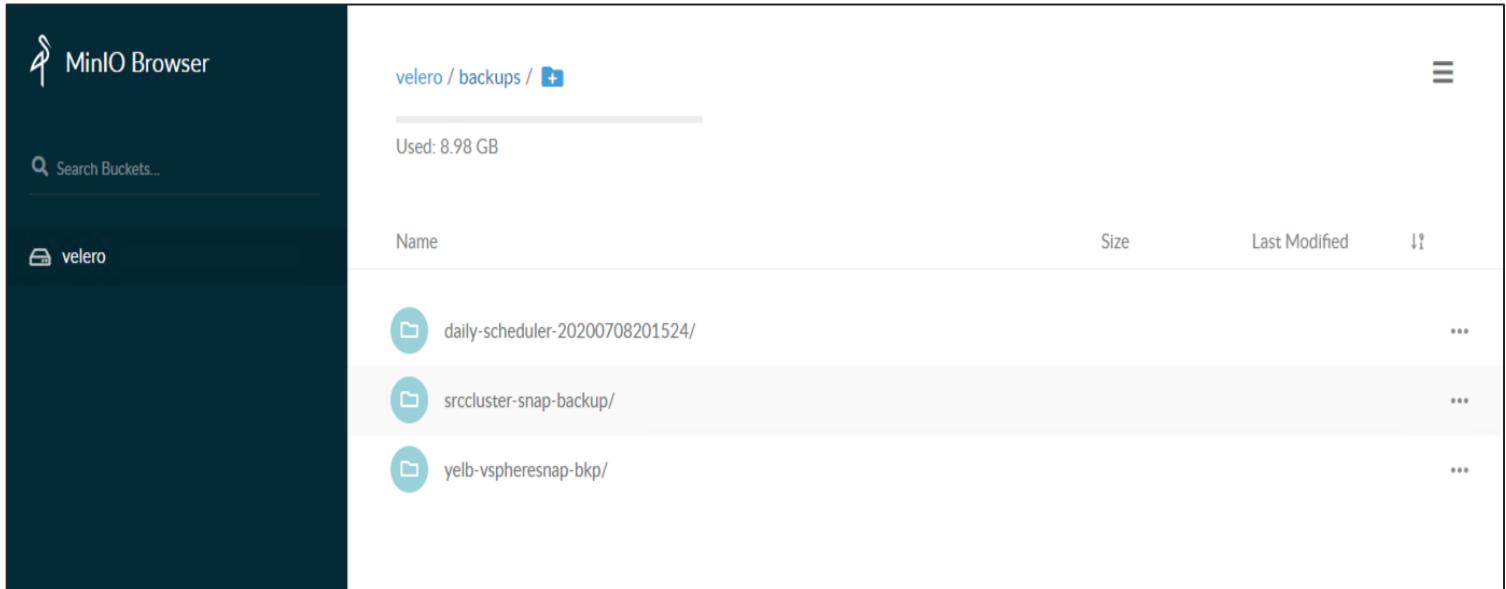
TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl -n velero get uploads
NAME                                AGE
upload-23500ea7-e487-4313-91d7-27a6f7297c4c  20m
upload-26662c39-642a-4179-af29-a110e1d0292c  20m
upload-2d9203ed-ac81-4e98-b468-15ca21658766  35m
upload-2dcc8590-d02b-4bf2-9222-f51562c971c5  44m
upload-75423580-9dcc-405d-8e34-7b324a226022  44m
upload-864510ed-73f1-4afd-ba23-5be409eb6b50  44m
upload-8671ebef-f703-4525-8797-81d3e0e0fd9f  44m
upload-92dd880c-dd17-49fc-92d8-9ff5c57d5503  85s
upload-98905351-7d62-4803-bf94-1c14dea33757  89s
upload-994551e3-7627-47d3-aba8-d8a93028ef97  80s
upload-c6d903d2-edaf-41b2-b6de-6f2b339bf009  35m
upload-df77bd51-56db-45e7-9ae6-225e8fb6a3f7  35m
upload-ee4c66b3-ab98-4ea0-90e1-b6ecc6ca67d5  96s
upload-fd1460d6-710a-4c0b-9b1b-422fa5373078  35m
```

Step 5: Check vCenter to make sure that none of the operations related to the backup are running

NOTE: Event if the velero backup status shows complete the backup would not have been complete. When the vsphere plugin is used there is a lot of operations that happen in the background. During the backup there are a number of operations that occur in Vsphere.

Step 6: Login to Minio and check if the scheduled backup has been created



The screenshot shows the MinIO Browser interface. On the left, a sidebar displays the 'MinIO Browser' logo and a search bar labeled 'Search Buckets...'. Below the search bar is a list of buckets, with 'velero' being the selected bucket, indicated by a blue icon. The main pane shows the contents of the 'velero' bucket under the 'velero / backups /' path. The interface includes a progress bar at the top stating 'Used: 8.98 GB'. A table lists three objects:

Name	Size	Last Modified	Actions
daily-scheduler-20200708201524/			...
srccluster-snap-backup/			...
yelb-vspheresnap-bkp/			...

Backup Type	Votes	Outback	Buca	IHop	Chipotle
Full Cluster	15	4 (27%)	3(20%)	5(33%)	3(20%)
Namespace	38	12(8%)	8(21%)	14(37%)	4(11%)
Scheduled backup	47	17(36%)	8(17%)	18(38%)	4(8.51%)

Check Velero documentation <https://velero.io/docs/v1.4/> for other options

Restore Backups

This section describes steps to restore a Velero backup to a target cluster. The steps give an overview of restoring a backup of a namespace and an entire clusterbackup. The restore procedure is the same if using a backup that used the restic plugin or the vsphere plugin. If the status of the restore indicates complete it does mean the restore is complete unlike the backup process.

Step 1: Get kube config for the source cluster

```
pks login -a <pks api> -u <user> -p <password> -k  
pks get-credentials <cluster>
```

Alternatively

```
pks get-kubeconfig <cluster> -a <pks api> -u <user> -p <password> -k
```

E.g.

```
pks login -a pks.corp.local -u riaz -p VMware1! -k  
pks get-credentials my-cluster  
pks get-kubeconfig my-cluster -a pks.corp.local -u riaz -p VMware1! -k
```

Step 2: Set kubectl context to the target cluster

```
kubectl config use-context <target-cluster>  
E.g.  
kubectl config use-context my-cluster
```

Step 3: Check all resources running on the target cluster

```
kubectl get ns
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get ns
NAME      STATUS   AGE
default   Active   33d
kube-node-lease Active  33d
kube-public Active  33d
kube-system Active  33d
pks-system Active  33d
spc       Terminating 32d
velero   Active  10d
```

NOTE: yelb, x1, y1 and z1 namespaces do not exist

Step 4: Before backing up to a cluster, make sure you have defined a storage class to be used by the stateful applications that is being restored. Create a storage-class on the target cluster with the following storage class definition.

Copy the contents of the file below to a file storage-class.yaml and create the storage class.

```
---
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  storageclass.kubernetes.io/is-default-class: "true"
  name: thin-disk
provisioner: kubernetes.io/vsphere-volume
parameters:
  diskformat: thin
  
```

kubectl apply -f storage-class.yaml

```
kubectl get sc
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get sc
NAME          PROVISIONER          AGE
thin-disk (default)  kubernetes.io/vsphere-volume  33d
```

NOTE: If setting up the storage class using the csi driver , follow the steps provided in the VMware Tanzu documentation to set up the CSI driver on the cluster you before creating the storage class.

<https://docs.pivotal.io/pks/1-7/vsphere-cns.html>

Storage class definition when using a CSI driver

```
---
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: csi-sc
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
  provisioner: csi.vsphere.vmware.com
parameters:
  datastoreurl: "ds:///vmfs/volumes/5cef81a9-a9328547-8d05-00505601dfda/"
```

```
kubectl get sc
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get sc
NAME          PROVISIONER          AGE
csi-sc (default)  csi.vsphere.vmware.com  20h
```

Step 5: Create a Volume snapshot location. (If using the velero vsphere plugin)

NOTE: If a snapshot location has already been created this step is optional

```
cd ~/velero/velero-v1.4.0-linux-amd64  
./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero snapshot-location create vsphere-snap-loc --provider velero.io/vsphere  
Snapshot volume location "vsphere-snap-loc" configured successfully.
```

Restore Namespace Backup

Step 6: Restore the yelb namespace from the yelbbkup created in the previous step.

```
cd ~/velero/velero-v1.4.0-linux-amd64
./velero restore create --from-backup yelbbkup
```

In the above example for vsphere plugin the namespace backup was named yelb-vspheresnap-bkp

```
./velero restore create --from-backup yelb-vspheresnap-bkp
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero restore create --from-backup yelbbkup
Restore request "yelbbkup-20200609031343" submitted successfully.
Run `velero restore describe yelbbkup-20200609031343` or `velero restore logs yelbbkup-20200609031343` for more details.
```

Step 7: Check the status of the restore in the cluster. The yelb namespace should be created and the pods should be up and running. Make sure that the pv is also created and bound:

```
kubectl get ns
kubectl get po -n yelb
kubectl get pvc -n yelb
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get ns
NAME      STATUS    AGE
default   Active   33d
kube-node-lease  Active   33d
kube-public  Active   33d
kube-system  Active   33d
pks-system  Active   33d
spc        Terminating 32d
velero     Active   10d
yelb       Active   4m1s
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get po -n yelb
NAME          READY   STATUS    RESTARTS   AGE
redis-server-0 1/1     Running   0          4m10s
yelb-appserver-696b9668c4-19ds5 1/1     Running   0          4m10s
yelb-db-0      1/1     Running   0          4m9s
yelb-ui-6665575695-58dfk     1/1     Running   0          4m8s
```

kubectl get pvc -n yelb						
NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
db-pv-claim	Bound	pvc-309cfeda-c316-4812-ac1f-1c4a300a487a	5Gi	RWO	thin-disk	4m29s
redis-pv-claim	Bound	pvc-e5c5f9fe-1e69-4dd6-a4b4-81a01aacf6f2	2Gi	RWO	thin-disk	4m29s

Note: Only the yelb namespace and its resources have been restored

Step 8: Get the external ip of yelb-ui, point the browser to it and make sure all the data is visible in the application, and the application is reachable. Compare the voting data recorded in the table before the namespace backup was taken.

```
kubectl get svc -n yelb
```

kubectl get svc -n yelb					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
redis-server	ClusterIP	10.100.200.135	<none>	6379/TCP	47s
yelb-appserver	ClusterIP	10.100.200.52	<none>	4567/TCP	47s
yelb-db	ClusterIP	10.100.200.247	<none>	5432/TCP	47s
yelb-ui	LoadBalancer	10.100.200.146	10.40.14.48	80:32015/TCP	47s

TKG Backup and Restore Clusters

10.40.14.48

HOL Admin RegionA NSX Manager Ops Manager Harbor vRealize Log Insight Concourse Kubernetes Dashboard Lab Guides EPMC

Yelb

Welcome to Yelb, the only hub for healthy food recommendations!



IHOP

Pancakes, for a powerful start!

[VOTE](#)



Chipotle

Burritos, for a mid-day break!

[VOTE](#)



Outback

Blooming onion, what else?

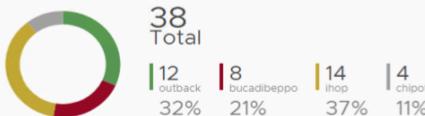
[VOTE](#)



Buca di Beppo

Lasagne, a great taste directly from Italy

[VOTE](#)



38 Total

Category	Count	Percentage
outback	12	32%
bucadibeppo	8	21%
ihop	14	37%
chipotle	4	11%

Number of page views so far: 1

System serving the request: [yelb-appserver-696b9668c4-nblp4](#)

Step 9: Delete the yelb namespace which will delete the application and the PV:

```
kubectl delete ns yelb
```

```
kubectl get ns
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl delete ns yelb
namespace "yelb" deleted
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get ns
NAME      STATUS   AGE
default   Active  33d
kube-node-lease  Active  33d
kube-public  Active  33d
kube-system  Active  33d
pks-system  Active  33d
spc        Terminating  32d
velero    Active  10d
```

Restore Cluster Backup

Step 10: Restore the back of all resources from the source cluster to the target cluster:

```
cd ~/velero/velero-v1.4.0-linux-amd64
./velero restore create --from-backup sourceclusterbk
```

In the above example for vsphere plugin the namespace backup was named yelb-vspheresnap-bkp

```
./velero restore create --from-backup srcluster-snap-backup
```

TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero restore create --from-backup sourceclusterbk
Restore request "sourceclusterbk-20200601213029" submitted successfully.
Run `velero restore describe sourceclusterbk-20200601213029` or `velero restore logs sourceclusterbk-20200601213029` for more details.
```

Step 11: Monitor the resources created in the target cluster. The yelb , x1, y1 and z1 namespaces should be created. Pods,pv's, deployments and services should also be created.

```
kubectl get ns  
kubectl get po --all-namespaces  
kubectl get pvc --all-namespaces  
kubectl get svc --all-namespaces
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get ns  
NAME          STATUS    AGE  
default        Active   33d  
kube-node-lease  Active   33d  
kube-public    Active   33d  
kube-system    Active   33d  
pks-system     Active   33d  
spc            Terminating 32d  
velero         Active   10d  
x1             Active   104s  
y1             Active   102s  
yelb           Active   106s  
z1             Active   104s
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get pvc --all-namespaces  
NAMESPACE  NAME      STATUS    VOLUME                                     CAPACITY  ACCESS MODES  STORAGECLASS  AGE  
spc        database-server  Terminating  pvc-7c5d2222-389d-4dfd-9187-a52d4db37a06  8Gi       RWO          thin-disk      32d  
yelb       db-pv-claim   Bound      pvc-74a65359-4a58-4087-a79e-01d67598197e  5Gi       RWO          thin-disk      4m35s  
yelb       redis-pv-claim Bound      pvc-bcfdb7eb-78e1-400d-b58c-3e954b8dd9ad  2Gi       RWO          thin-disk      4m35s
```

TKG Backup and Restore Clusters

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	33d
kube-system	kube-dns	ClusterIP	10.100.200.2	<none>	53/UDP, 53/TCP	33d
kube-system	kubernetes-dashboard	NodePort	10.100.200.95	<none>	443:31379/TCP	33d
kube-system	metrics-server	ClusterIP	10.100.200.6	<none>	443/TCP	33d
pks-system	fluent-bit	ClusterIP	10.100.200.80	<none>	24224/TCP	33d
pks-system	validator	ClusterIP	10.100.200.207	<none>	443/TCP	33d
x1	service-a-lb	LoadBalancer	10.100.200.147	10.40.14.138	80:32641/TCP	9m11s
x1	svc-service-a	ClusterIP	10.100.200.22	<none>	80/TCP	9m11s
y1	svc-service-b	ClusterIP	10.100.200.138	<none>	80/TCP	9m11s
yelb	redis-server	ClusterIP	10.100.200.4	<none>	6379/TCP	9m11s
yelb	yelb-appserver	ClusterIP	10.100.200.5	<none>	4567/TCP	9m11s
yelb	yelb-db	ClusterIP	10.100.200.133	<none>	5432/TCP	9m11s
yelb	yelb-ui	LoadBalancer	10.100.200.33	10.40.14.139	80:30859/TCP	9m11s
z1	svc-service-c	ClusterIP	10.100.200.109	<none>	80/TCP	9m13s
z1	svc-service-d	ClusterIP	10.100.200.112	<none>	80/TCP	9m13s

Step 12: Get the external ip for the yelb-ui app

```
kubectl get svc -n yelb
```

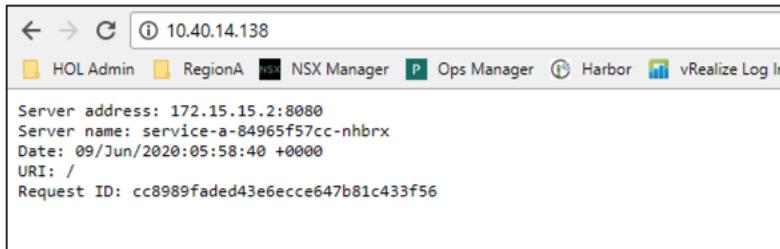
```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get svc -n yelb
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP     PORT(S)        AGE
redis-server   ClusterIP 10.100.200.4 <none>        6379/TCP      9m46s
Yelb-appserver ClusterIP 10.100.200.5 <none>        4567/TCP      9m46s
yelb-db        ClusterIP 10.100.200.133 <none>       5432/TCP      9m46s
Yelb-ui        LoadBalancer 10.100.200.33  10.40.14.139  80:30859/TCP  9m46s
```

Step 13: Point the browser to it and make sure all the data is visible in the application and the application is reachable. Compare the voting data recorded in the table before the cluster backup was taken:

The screenshot shows the Yelb application running in a browser. The URL is 10.40.14.139. The page displays four food options: IHOP, Chipotle, Outback, and Buca di Beppo. Each card includes a logo, a short description, and a blue 'VOTE' button. Below the cards is a pie chart titled '15 Total' showing the distribution of votes: Outback (4, 27%), Buca di Beppo (3, 20%), IHOP (5, 33%), and Chipotle (3, 20%). At the bottom left is a box for 'Number of page views so far: 1'. At the bottom right is a box stating 'System serving the request: yelb-appserver-696b9668c4-nblp4'.

Step 14: The x1, y1 and z1 namespaces run nginx pods with busybox. The nginx pod running in namespace x1 is exposed as a load balancer. Get the external ip to this pod and point the browser to access the service:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service-a-lb	LoadBalancer	10.100.200.179	10.40.14.138	80:32017/TCP	4m17s
svc-service-a	ClusterIP	10.100.200.97	<none>	80/TCP	4m16s



Note: With the cluster restore all resources , including namespaces, pv's etc are restored

From a cluster backup you could selectively restore the resources that are required to migrate as well.

Restore a Point in Time Backup (Scheduled Backup)

Step 14: Delete the yelb, x1 , y1 and z1 namespaces which will delete the application and the associated PV's and other resources:

```
kubectl delete ns yelb  
kubectl delete ns x1  
kubectl delete ns y1  
kubectl delete ns z1
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl delete ns yelb  
namespace "yelb" deleted  
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl delete ns x1  
namespace "x1" deleted  
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl delete ns y1  
namespace "y1" deleted  
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl delete ns z1  
namespace "z1" deleted
```

TKG Backup and Restore Clusters

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get ns
NAME        STATUS   AGE
default     Active   34d
kube-node-lease Active   34d
kube-public  Active   34d
kube-system  Active   34d
pks-system   Active   34d
spc         Terminating   32d
velero      Active   30m
```

Step 15: Login to Minio, and copy the name of the backup that was created when a backup scheduler was created. Eg. daily-scheduler-20200609044342:

Step 15: Restore the point in time backup to the cluster:

```
cd ~/velero/velero-v1.4.0-linux-amd64
./velero restore create --from-backup daily-scheduler-20200609044342
```

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ ./velero restore create --from-backup daily-scheduler-20200609044342
Restore request "daily-scheduler-20200609044342-20200609061417" submitted successfully.
Run `velero restore describe daily-scheduler-20200609044342-20200609061417` or `velero restore logs daily-scheduler-20200609044342-20200609061417` for more details.
```

Step 16: Check the status of the restore:

TKG Backup and Restore Clusters

./velero restore describe daily-scheduler-20200609044342-20200609061417

```
Phase: Completed
Warnings:
Velero: <none>
Cluster: could not restore, namespaces.velero.com "8b177990-3517-4e8d-83af-02b5ff53e937" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, namespaces.velero.com "selection-lock-pke-522f20e3-c3b0-4d8a-81ee-0b0faefc2e46" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, veleroconfigurations.admissionregistration.k8s.io "validator-pkapi.io" already exists. Warning: the in-cluster version is different than the backed-up version.
Namespaces:
default:
        could not restore, pods "busybox-7b87495f89-1ge4f" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "kube-public" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "kube-system" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "kubernetes" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "kubernetes-dashboard-certs" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "metrics-server-certs" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "node-defaults" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "node-labels" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, endpoints "kube-controller-manager" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, endpoints "kube-dns" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, endpoints "kube-dns" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, endpoints "kube-dns" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, endpoints "kube-dns" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, services "kube-dns" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, services "metrics-server" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "fluent-bit" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "fluent-bit" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, secrets "validator" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, configmaps "telemetry-config-map" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, configmaps "telemetry-config-map" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, endpoints "validator" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, services "validator" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, services "validator" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, podsvolumebackups.velero.io "sourceclusters-4fd46" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, podsvolumebackups.velero.io "yelb-backup-3fdab" already exists. Warning: the in-cluster version is different than the backed-up version.
        could not restore, podsvolumebackups.velero.io "yelb-backup-3fdab" already exists. Warning: the in-cluster version is different than the backed-up version.

Backup: daily-scheduler-20200609044342
Namespaces: Included: all namespaces found in the backup
            <none>
Resources: Included:
            nodes, events, events.events.k8s.io, backups.velero.io, restores.velero.io, resticrepositories.velero.io
Cluster-scoped: auto
Namespace mappings: <none>
Label selector: <none>
Restore PVs: auto
Restic Restores (specify --details for more information):
Completed: 2
```

Step 17: Monitor the resources created in the target cluster. The yelb , x1, y1 and z1 namespaces should be created. Pods,pv's, deployments and services should also be created.

kubectl get ns

kubectl get po --all-namespaces

kubectl get pvc --all-namespaces

kubectl get svc --all-namespaces

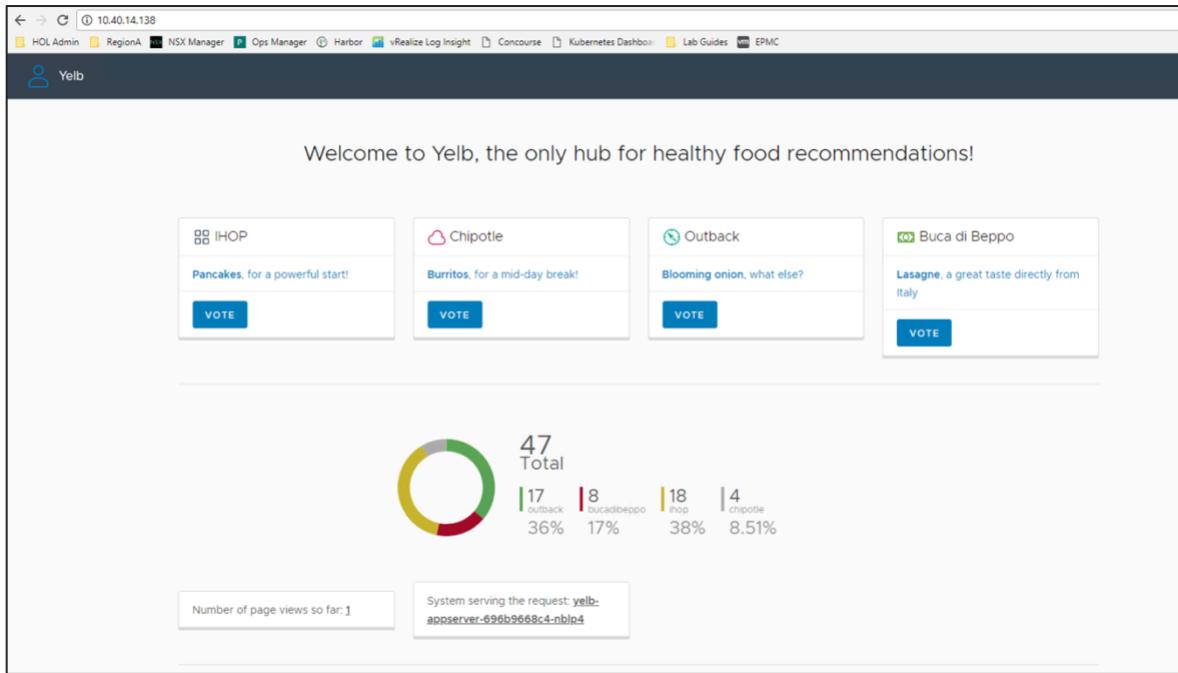
Step 18: Get the external ip for the yelb-ui app

kubectl get svc -n yelb

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get svc -n yelb
NAME         TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)        AGE
redis-server ClusterIP  10.100.200.110 <none>        6379/TCP     3m8s
yelb-appserver ClusterIP  10.100.200.206 <none>        4567/TCP     3m8s
yelb-db       ClusterIP  10.100.200.157 <none>        5432/TCP     3m7s
yelb-ui       LoadBalancer 10.100.200.75  10.40.14.138  80:31757/TCP  3m7s
```

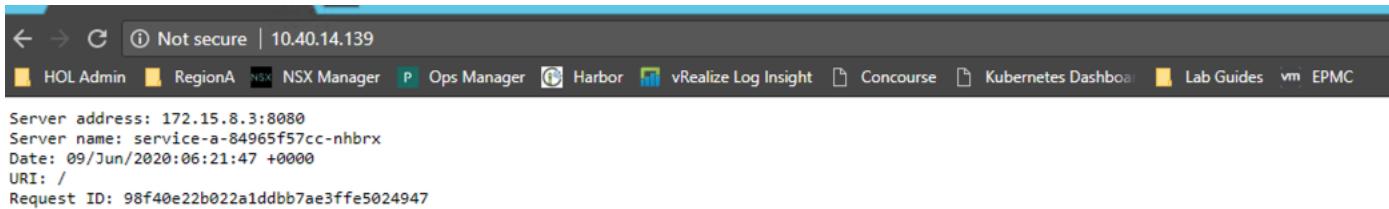
TKG Backup and Restore Clusters

Step 19: Point the browser to it, make sure all the data is visible in the application, and the application is reachable. Compare the voting data recorded in the table before the cluster backup was taken.



Step 20: The x1, y1 and z1 namespaces run nginx pods with busybox. The nginx pod running in namespace x1 is exposed as a load balancer. Get the external ip to this pod and point the browser to access the service.

```
ubuntu@cli-vm:~/velero/velero-v1.4.0-linux-amd64$ kubectl get svc -n x1
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service-a-lb   LoadBalancer   10.100.200.216   10.40.14.139   80:31705/TCP   5m28s
svc-service-a  ClusterIP    10.100.200.126   <none>        80/TCP       5m28s
```



Note: With the cluster restore all resources , including namespaces, pv's etc are restored from a backup in

time.

Cleanup

Step 1: Get kube config for the source cluster

```
pks login -a <pks api> -u <user> -p <password> -k  
pks get-credentials <cluster>
```

Alternatively

```
pks get-kubeconfig <cluster> -a <pks api> -u <user> -p <password> -k
```

E.g.

```
pks login -a pks.corp.local -u riaz -p VMware1! -k  
pks get-credentials ci-cluster  
pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k
```

```
ubuntu@cli-vm:~/velero$ pks get-kubeconfig ci-cluster -a pks.corp.local -u riaz -p VMware1! -k  
Fetching kubeconfig for cluster ci-cluster and user riaz.  
You can now use the kubeconfig for user riaz:  
$kubectl config use-context ci-cluster
```

Step 2: Set kubectl context to the source cluster

```
kubectl config use-context <source-cluster>
```

E.g.

```
kubectl config use-context ci-cluster
```

Step 3: Get all backup's taken from of a specific cluster

```
./velero backup get
```

Step 4: To delete a specific backup

```
./velero backup delete <backupname>
```

E.g.

```
./velero backup delete sourceclusterbk
```

Step 4: To delete all backups

```
./velero backup delete --all
```

Step 5: To delete a backup schedule

```
./velero schedule delete <schedule-name>
```

Conclusion

We hope this document was useful. As you try these configuration steps, please provide any feedback or questions in the comments section for this document on code.vmware.com. Also, do let us know if you have any suggestions or if you would like to see guidance on other topics.



VMware, Inc. 3401 Hillview Avenue Palo Alto CA 94304 USA Tel 877-486-9273 Fax 650-427-5001 www.vmware.com.
Copyright © 2019 VMware, Inc. All rights reserved. This product is protected by U.S. and international copyright and intellectual property laws. VMware products are covered by one or more patents listed at vmware.com/go/patents. VMware is a registered trademark or trademark of VMware, Inc. and its subsidiaries in the United States and other jurisdictions. All other marks and names mentioned herein may be trademarks of their respective companies. Item No: vmw-wp-tech-temp-word-102-proof 5/19